



Intro to Rust Lang Ownership (Part 1)

Welcome back!

- Homework 1 due in 2 weeks
- You can use 7 late days over the whole semester
- If you spent over an hour on the assignment, please let us know!

Today: Ownership

- Ownership
 - Motivation: Strings
 - Moving, `clone`, and `copy`
- References and Borrowing
- Slices and Owned Types

Ownership

Ownership

From the official Rust Lang [book](#):

Ownership is Rust's most unique feature and has deep implications for the rest of the language. It enables Rust to make memory safety guarantees without needing a garbage collector, so it's important to understand how ownership works.

- Today we'll introduce *Ownership*, as well as several related features

What is Ownership?

Ownership is a set of rules that govern how a Rust program manages memory.

- Some languages have garbage collection to manage memory
- Other languages require you to explicitly allocate and free memory
- Rust has a third approach:
 - Manage resources via a set of *rules*

Ownership Rules

- Each value in Rust has an *owner*
- There can only be one owner at a time
- When the owner goes out of scope, the value will be *dropped*

Ownership Rules

For now, assume the following:

- Variables **own** values
- Variables can store their values in two places:
 - Stored on the stack
 - Stored somewhere else...

Motivation: Strings

- Last week: Scalar Data Types
 - Values live on the stack
- This week: `String s`
 - Values live somewhere else...

String Literals

Every time you see a text like "Hello, World!" surrounded by double quotes, that is a *string literal*.

```
fn main() {  
    println!("Hello, world!");           // Print a string literal  
  
    let s = "Ferris is our friend";    // Store another string literal  
}
```

- String literals are stored in a read-only section of the program binary
 - In other words, these strings literals are known at *compile-time*

Python Strings

Suppose we wanted store a user's input. Python can do this easily!

```
username1 = input("Enter a short name: ")      # Could be "Bob"
username2 = input("Enter a long name: ")        # Could be "Bartholomew"

# Python strings can be any length!
assert len(username1) == 3 and len(username2) == 11
```

- In Python, we don't have to know the length of the string beforehand
- How would we do this in Rust?

Problem: String Literals are Immutable

String literals in Rust must be known at compile-time, so we cannot use them for this type of program.

```
fn main() {  
    let username1: <???> = ask_for_user_input();  
    let username2: <???> = ask_for_user_input();  
}
```

- We don't know size of `username` at compile time
- These strings must be *dynamically sized*

The `String` Type

In addition to string literals, Rust has another string type called `String`.

- `String` manages data allocated on the *heap*
- We use `String` for managing string data where we do not know the size of the string at compile-time

String Example

You can create a `String` from a string literal using `String::from()`.

```
let s = String::from("hello");
```

String Example

This kind of string *can* be mutated:

```
let mut s = String::from("hello");

s.push_str(", world!"); // push_str() appends a literal to a String

println!("{}", s); // This will print `hello, world!`
```

Problem: When is `String` valid?

- String literals are hardcoded into the executable
 - Always valid ✓
- On the other hand, `String`s are dynamically allocated on the *heap*
 - Must request memory from the allocator at *runtime*
 - Must return the memory when we're done using it
 - Who's responsible for this?

Python and Java's Proposal: Garbage Collection

In Python and Java, the *runtime* is responsible for managing memory.

- The runtime is a system that provides services while the program runs
 - Among these services is the garbage collector
 - The garbage collector finds unused memory and cleans it up
 - Runtime processes can be inefficient!

C's Proposal: Manual Memory Management

In C, the *programmer* is responsible for returning memory.

- `malloc` : request memory from allocator
- `free` : return memory to allocator

```
int main() {
    char *s = (char *)malloc(sizeof("hello")); // Allocate memory for `s`
    strcpy(s, "hello");                      // Set `s` to "hello"
    free(s);                                // Done with `s`, free it!
}
```

C's Proposal: Manual Memory Management

However, the programmer must pair exactly one `malloc` with exactly one `free`.

- If we forget to `free`, we leak memory
- If we `free` too early, we have an invalid variable
- If we `free` twice, that's a "double free" bug
- Undefined behavior!!! 

C's Proposal: Memory Footgun

Using `malloc` and `free` can lead to all sorts of undefined behavior.

- Unless you are the perfect developer...
- Who *never* writes a bug...
- You're bound to shoot yourself in the foot!

Rust's Proposal: Compile-Time Memory Safety

In Rust, the *compiler* is responsible for returning memory.

- Compiler marks places to return memory
- It would be great if the compiler knew:
 - When the variable needs memory
 - When the memory is no longer needed
- Idea: **What if we tied heap allocations to the scope of a variable?**

Rust's Proposal: Compile-Time Memory Safety

Every variable has a *scope*.

```
{  
    let s = String::from("hello"); // s comes into scope  
} // s goes out of scope
```

- There are two important points here:
 - When `s` comes *into* scope, it is valid
 - When `s` goes *out of* scope, it is invalid

Rust's Proposal: Compile-Time Memory Safety

Memory is returned once the variable that owns it goes out of scope.

```
{  
    let s = String::from("hello"); // s comes into scope  
} // s goes out of scope
```

- When `s` comes into scope, it gets memory from the allocator
- When `s` goes out of scope, its memory is freed
 - Rust automatically calls the function `drop` on `s` when we reach the closing bracket

Example: String "copying"

```
let s1 = String::from("hello");
let s2 = s1;
```

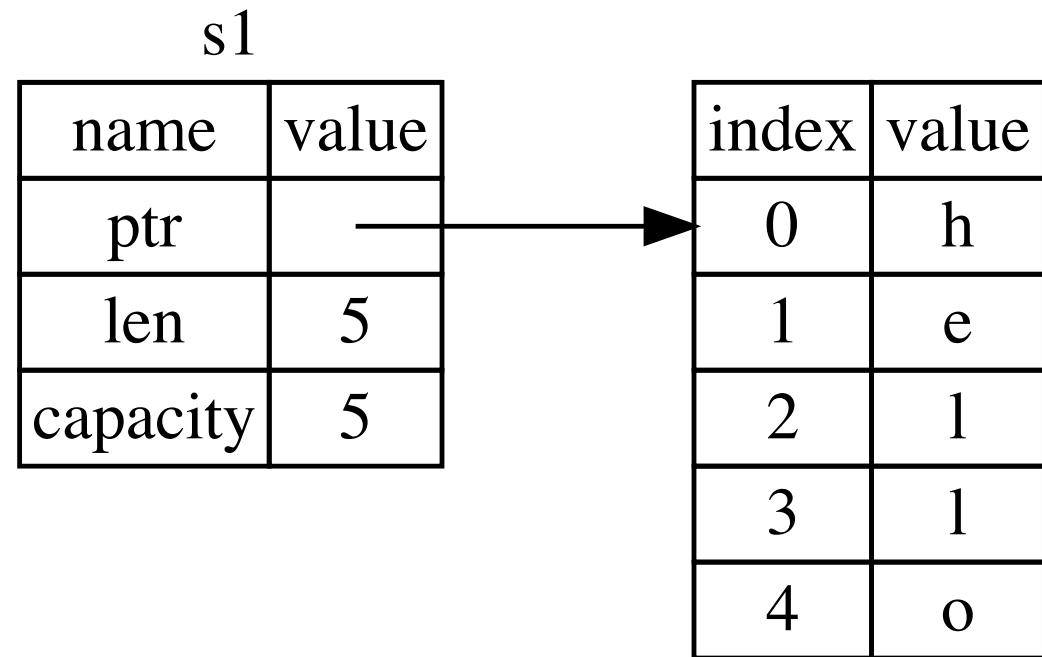
What is this code doing?

- Bind the `String` containing "hello" to `s1`
- Now what?
 - Do we make a copy of the `String` ?
 - What does a copy actually mean in this case?

String Data Layout

```
let s1 = String::from("hello");
```

- A `String` is made up of 3 fields:
 - A pointer to the characters somewhere in memory
 - A length
 - A capacity
- Left diagram is on the stack
- Right diagram is on the heap



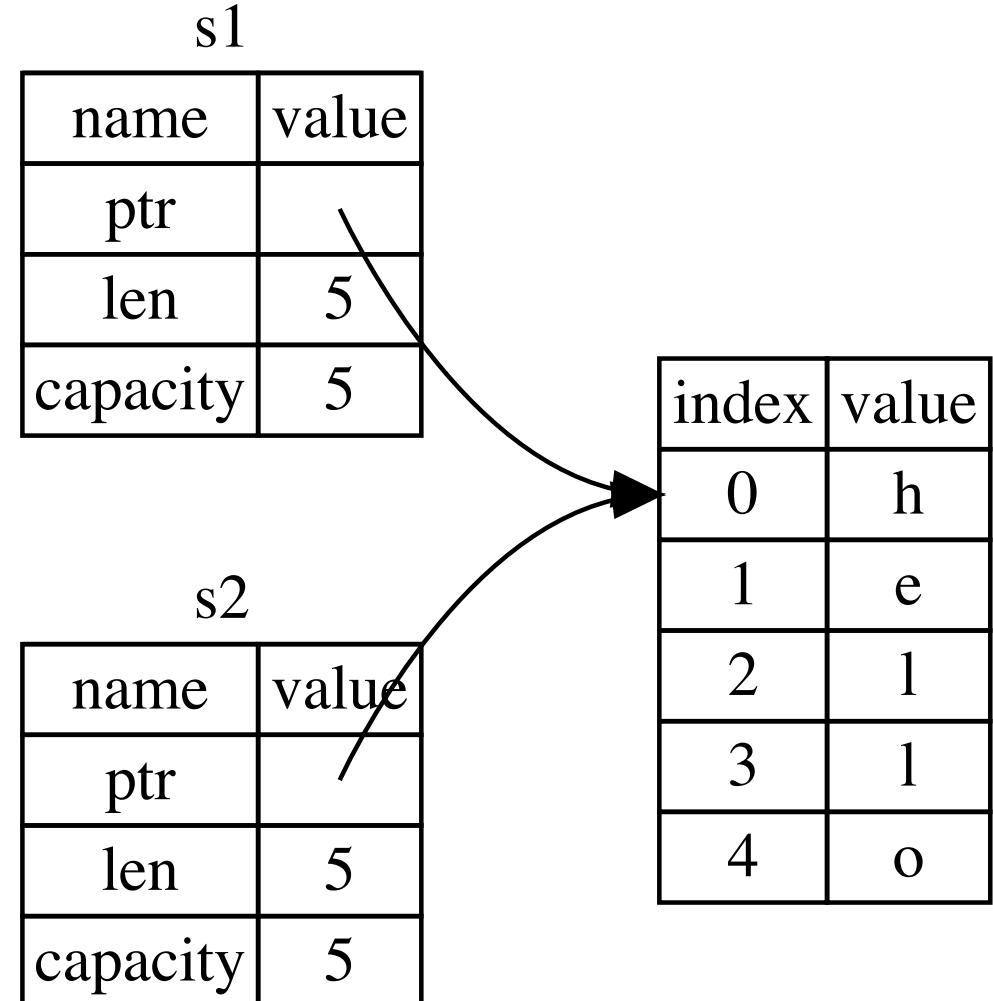
Pointer Aliasing



```
let s1 = String::from("hello");
let s2 = s1;
```

One way to handle this case is:

- When we assign `s1` to `s2`, only the stack data is copied
- We do *not* create a copy of the contents of the `String`
- Also known as a "shallow copy" in some languages

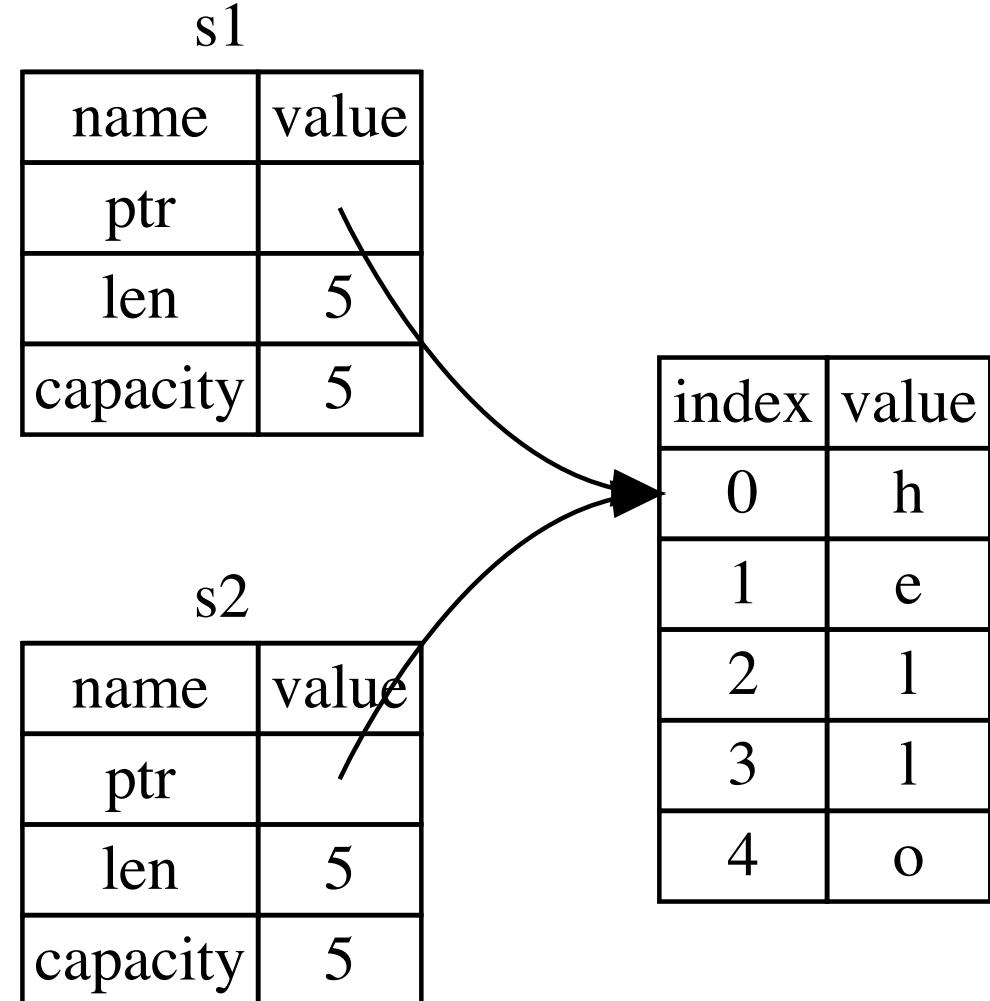


Pointer Aliasing 💀

```
let s1 = String::from("hello");
let s2 = s1;
```

Suppose Rust handled this case with a shallow copy.

- Following Rust's scope rules, what would happen if we tried to drop both `s1` and `s2`?
 - Double free! 🪢
- How can this be prevented?

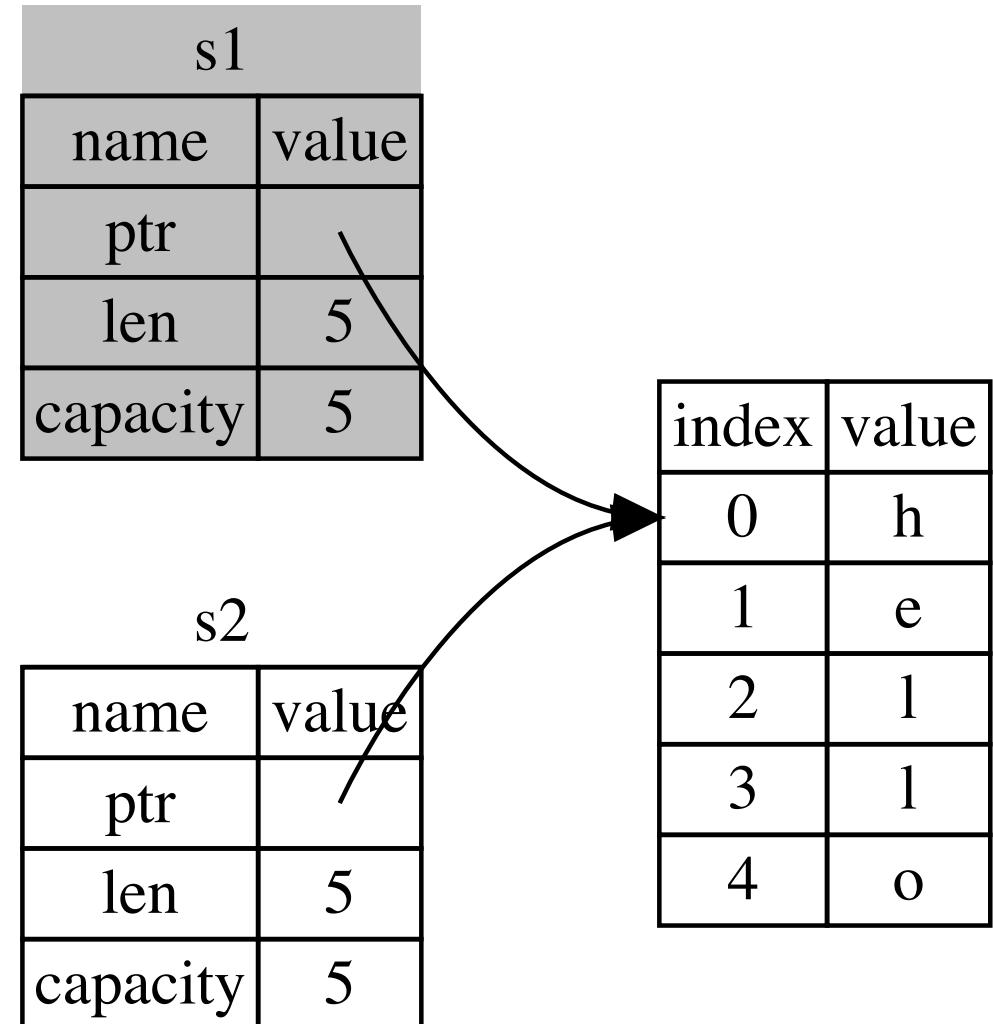


Enforcing Single Ownership

To ensure memory safety, after the second line, `s1` is no longer valid.

```
let s1 = String::from("hello");
let s2 = s1; // s1 is no longer valid
```

- *Grayed out portion is no longer accessible to the program*



What happens if we try to use `s1` after it is invalid?

```
let s1 = String::from("hello");
let s2 = s1;
println!("{} , world!", s1);
```

```
error[E0382]: borrow of moved value: `s1`
|
2 |     let s1 = String::from("hello");
|         -- move occurs because `s1` has type `String`,
|             which does not implement the `Copy` trait
3 |     let s2 = s1;
|             -- value moved here
4 |
5 |     println!("{} , world!", s1);
|                         ^^ value borrowed here after move
```

Move Semantics

```
let s1 = String::from("hello");      // Create `s1`  
let s2 = s1;                      // Move `s1` into `s2`  
// println!("{} , world!", s1);    // `s1` is now invalid!
```

- Rust calls this shallow copy plus invalidation a *move*
- We *moved* `s1` into `s2`
 - Hence `s1` can no longer be accessed

Moving vs Cloning

```
let s1 = String::from("hello");
let s2 = s1;
```

- What if we *wanted* to copy all of the data?
 - Known as deep copying or cloning
- Making Rust implicitly deep copy all data would solve the problem
 - But it would get expensive, quickly

Clone

If we do want to deep copy, we can use a method called `clone`.

```
let s1 = String::from("hello");
let s2 = s1.clone();

println!("s1 = {}, s2 = {}", s1, s2);
```

```
s1 = hello, s2 = hello
```

- This copies *all* of the data contained in `s1`, both on the heap and the stack

Clone

```
let s1 = String::from("hello");
let s2 = s1.clone();

println!("s1 = {}, s2 = {}", s1, s2);
```

```
s1 = hello, s2 = hello
```

- In Rust, cloning must be explicitly performed by the programmer
 - This is very intentional, to avoid accidental performance overhead
- We'll talk more about methods next week!

What About Integers?

Based on the ownership rules we have stated, this code should not work.

```
let x = 5;  
let y = x;  
  
println!("x = {}, y = {}", x, y);
```

```
x = 5, y = 5
```

- `x` is still valid, but it looks like we moved it into `y`
- Didn't we just say that this wasn't allowed?!

Copy

```
let x = 5;  
let y = x;
```

- Types such as integers have a size known at compile time
- Data is stored either in registers or on the stack
- Copies of integers are very quick to make
- There is no difference between a shallow copy and a deep copy here
 - So why not clone implicitly?



Copy

Certain types are annotated with a `Copy` trait, which allows implicit copying instead of moving.

Types that are `Copy` :

- All numeric types, including integers (`i32`) and floating points (`f64`)
- The boolean type (`bool`)
- The character type (`char`)
- Tuples, if they only contain types that are `Copy`
 - `(i32, i32)` is `Copy`, but `(i32, String)` is not



Ownership and Functions

Passing a variable to a function behaves just as assignment does.

Passing a `String`:

```
fn main() {
    let s = String::from("hello");
    takes_ownership(s);
} // Because `s`'s value was moved, `s` is not dropped

                    // `some_string` comes into scope
fn takes_ownership(some_string: String) {
    println!("{} is mine now!", some_string);
} // `some_string` goes out of scope and `drop` is called.
// The backing memory is freed.
```

Ownership and Functions

What if we tried to use a value after a function takes ownership of it?

```
let s = String::from("hello");
takes_ownership(s);
println!("{} is invalid now!", s);
```

```
error[E0382]: borrow of moved value: `s`
--> src/main.rs:4:36
|
2 |     let s = String::from("hello");
|         - move occurs because `s` has type `String`,
|           which does not implement the `Copy` trait
3 |     takes_ownership(s);
|                 - value moved here
4 |     println!("{} is invalid now!", s);
|                           ^ value borrowed here after move
```

Ownership and Functions

Copy are copied directly into the function parameter:

```
fn main() {
    let x = 5;
    makes_copy(x);
    println!("Here is {} again!", x); // x is still valid!
}
```

```
fn makes_copy(some_integer: i32) {
    println!("{} just got copied", some_integer);
}
```

```
5 just got copied
Here is 5 again!
```

Return Values and Scope

Returning values can also transfer ownership back to the caller.

```
fn main() {
    let s1 = gives_ownership();
    println!("{}", s1); // s1 is valid, we have taken ownership!
}

fn gives_ownership() -> String {
    let some_string = String::from("inside `gives_ownership`");
    some_string // `some_string` is returned and is moved out to the
                // calling function
}
```

inside `gives_ownership`

Return Values and Scope

Here is another example where a function takes ownership and gives it back:

```
fn main() {
    let s2 = String::from("hello");
    let s3 = takes_and_gives_back(s2);
    println!("{}", s3);
} // Here, `s3` goes out of scope and is dropped.
// `s2` was moved, so nothing happens to `s2`.

fn takes_and_gives_back(a_string: String) -> String {
    a_string // a_string is returned and
             // moves out to the calling function
}
```

Recap: Ownership

- Ownership rules:
 - Each value in Rust has an *owner*
 - There can only be one owner at a time
 - When the owner goes out of scope, the value will be *dropped*
- With just ownership, we can either move, copy, or clone
 - Moving and copying has no overhead
 - Cloning is expensive

Moving is somewhat tedious

```
fn main() {
    let s1 = String::from("hello");
    let (s2, len) = calculate_length(s1);
    println!("The length of '{}' is {}.", s2, len);
}
fn calculate_length(s: String) -> (String, usize) {
    let length = s.len();
    (s, length)
}
```

- If we want to give a function some data, it seems we need to *move* the data into the function
- To get it back, it seems we need to also return the data back every time
- *What if we want to let a function use a value but not take ownership?*

References and Borrowing

References

- Moving into and returning data from a function is a lot of work
- Rust has a feature specifically for using a value without transferring ownership called *references*
- We can share memory using these *references*

Reference with &

Instead of moving a value into a function, we can provide a *reference* to the value. We use & to define a reference to a value.

```
fn main() {
    let s1 = String::from("hello");
    let len = calculate_length(&s1);
    println!("The length of '{}' is {}.", s1, len);
}
fn calculate_length(borrowed: &String) -> usize {
    borrowed.len()
}
```

- The &s1 syntax lets us create a variable that *refers* to the value of s1
- &String means the type of the argument is a reference to a String

References as Function Arguments

```
// `borrowed` is a reference to a String
fn calculate_length(borrowed: &String) -> usize {
    borrowed.len()
} // Here, `borrowed` goes out of scope
```

- `borrowed` is a reference to `s1` (i.e. `&s1`)
- We *do not own* `s1` with just a reference to it
- This means `s1` will *not* be dropped when `borrowed` goes out of scope
- We call holding a reference *borrowing*

Mutating a Reference

What if we want to modify the value of something we've borrowed through a reference?

```
fn main() {  
    let s = String::from("hello");  
  
    change(&s);  
}  
  
fn change(some_string: &String) {  
    some_string.push_str(", world");  
}
```



Modifying a Reference

We get an error if we try to modify a reference.

```
error[E0596]: cannot borrow `*some_string` as mutable,
               as it is behind a `&` reference
--> src/main.rs:8:5
 |
7 | fn change(some_string: &String) {
|                               ----- help: consider changing this
|                               to be a mutable reference: `&mut String`  

8 |     some_string.push_str(", world");
|     ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^ `some_string` is a `&` reference,
|                               so the data it refers to cannot
|                               be borrowed as mutable
```

- Just like variables, references are immutable by default

Mutable References

If we want to modify the value that we've borrowed, we must use a mutable reference, denoted `&mut val`.

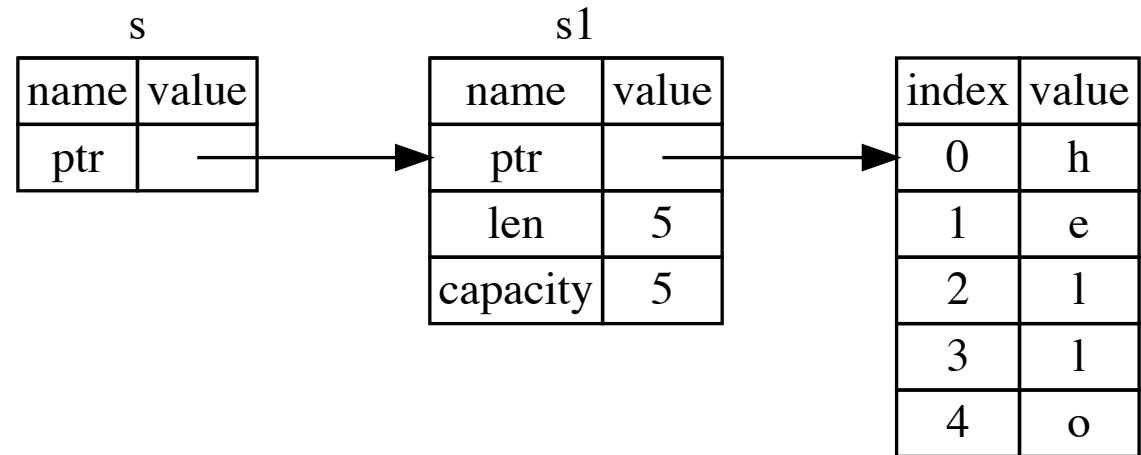
```
fn main() {
    let mut s = String::from("hello");

    change(&mut s);
}

fn change(some_string: &mut String) {
    some_string.push_str(", world");
}
```

Reference Data Layout

- In memory, references are just like pointers
- In practice, they have a couple of constraints that make them safer



Reference Constraints

- Mutable references must be exclusive
 - There can only be one mutable reference to a value at a time
- There can never be dangling references

Constraint: Mutable References are Exclusive

If you have a mutable reference to a value, you can have no other references to that value.

```
let mut s = String::from("hello");

let r1 = &mut s;
let r2 = &mut s;

println!("{} , {}", r1, r2);
```



Constraint: Mutable References are Exclusive

```
let mut s = String::from("hello");
let r1 = &mut s;
let r2 = &mut s;
println!("{} , {}", r1, r2);
```

```
error[E0499]: cannot borrow `s` as mutable more than once at a time
--> src/main.rs:5:14
|
4 |     let r1 = &mut s;
|           ----- first mutable borrow occurs here
5 |     let r2 = &mut s;
|           ^^^^^^ second mutable borrow occurs here
6 |
7 |     println!("{} , {}", r1, r2);
|           -- first borrow later used here
```

Constraint: Mutable References are Exclusive

- Most languages will let you mutate anything, whenever you want
- If data can be written to from multiple places, the value can become unpredictable
- Making mutable references exclusive can prevent data invalidation and data races at compile time!

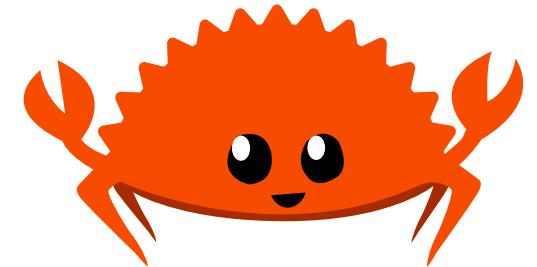
Multiple Mutable References

We are allowed to hold multiple mutable references, just not *simultaneously*.

```
let mut s = String::from("hello");

{
    let r1 = &mut s;
} // r1 goes out of scope here, so we can make
   // a new mutable reference with no problems

let r2 = &mut s;
```



- Notice that the scopes of these mutable references do not overlap

Mutable and Immutable References

We cannot have both an immutable and mutable reference to the same value.

```
let mut s = String::from("hello");

let r1 = &s; // no problem
let r2 = &s; // no problem
let r3 = &mut s; // BIG PROBLEM

println!("{} , {} , and {}", r1, r2, r3);
```



Mutable and Immutable References

```
error[E0502]: cannot borrow `s` as mutable because
                  it is also borrowed as immutable
--> src/main.rs:6:14
|
4 |     let r1 = &s; // no problem
   |             -- immutable borrow occurs here
5 |     let r2 = &s; // no problem
6 |     let r3 = &mut s; // BIG PROBLEM
   |             ^^^^^^ mutable borrow occurs here
7 |
8 |     println!("{} , {} , and {}", r1, r2, r3);
   |                         -- immutable borrow later used here
```

Mutable and Immutable References

Note that exclusivity rules only apply for references whose scopes overlap.

```
let mut s = String::from("hello");

let r1 = &s; // no problem
let r2 = &s; // no problem
println!("{} and {}", r1, r2);
// variables r1 and r2 will not be used after this point

let r3 = &mut s; // no problem
println!("{}", r3);
```

Mutable and Immutable References

```
let mut s = String::from("hello");
let r1 = &s; // no problem
let r2 = &s; // no problem
println!("{} and {}", r1, r2);

let r3 = &mut s; // no problem
println!("{}", r3);
```

- The scope of a reference starts when it is initialized
- The scope of a reference **ends at the last point it is used**
- The specific term for reference scopes are *lifetimes*
 - We'll talk about lifetimes in a future week!

Constraint: No Dangling References

The Rust compiler guarantees that references will never be invalid, which means it will not allow dangling references.

```
fn main() {
    let reference_to_nothing = dangle();
}

fn dangle() -> &String {
    let s = String::from("hello");

    &s
}
```



Constraint: No Dangling References

```
error[E0106]: missing lifetime specifier
--> src/main.rs:5:16
5 | fn dangle() -> &String {
    |          ^ expected named lifetime parameter
    |
    = help: this function's return type contains a borrowed value,
      but there is no value for it to be borrowed from
help: consider using the ``static`` lifetime
<-- snip -->
```

Focus on this line:

help: this function's return type contains a borrowed value, but there is no
value for it to be borrowed from



Reference Constraints

- Mutable references are exclusive:
 - At any given time, you can have either one mutable reference or any number of immutable references
 - A book being read by multiple people is fine
 - If multiple people write, they may overwrite each other's work
 - References are similar to Reader-Writer locks
- There can be no dangling references, references must always be valid

The Borrow Checker

The *Borrow Checker* enforces the ownership and borrowing rules by checking:

- That all variables are initialized before they are used
- That you can't move the same value twice
- That you can't move a value while it is borrowed
- That you can't access a place while it is mutably borrowed (except through the mutable reference)
- That you can't mutate a place while it is immutably borrowed
- and more...

Slices

Slices

- *Slices* let you reference a contiguous sequence of elements in a collection rather than the whole collection
- A slice is similar to a reference, so it does not have ownership

Slices

Suppose we want to write this function:

```
fn first_word(s: &String) -> ?
```

- Find the first space and return all the characters before it
- What type should we return?

String Slices

A *string slice* is sometimes a reference to part of a `String`, and it looks like this:

```
let s = String::from("hello world");
```

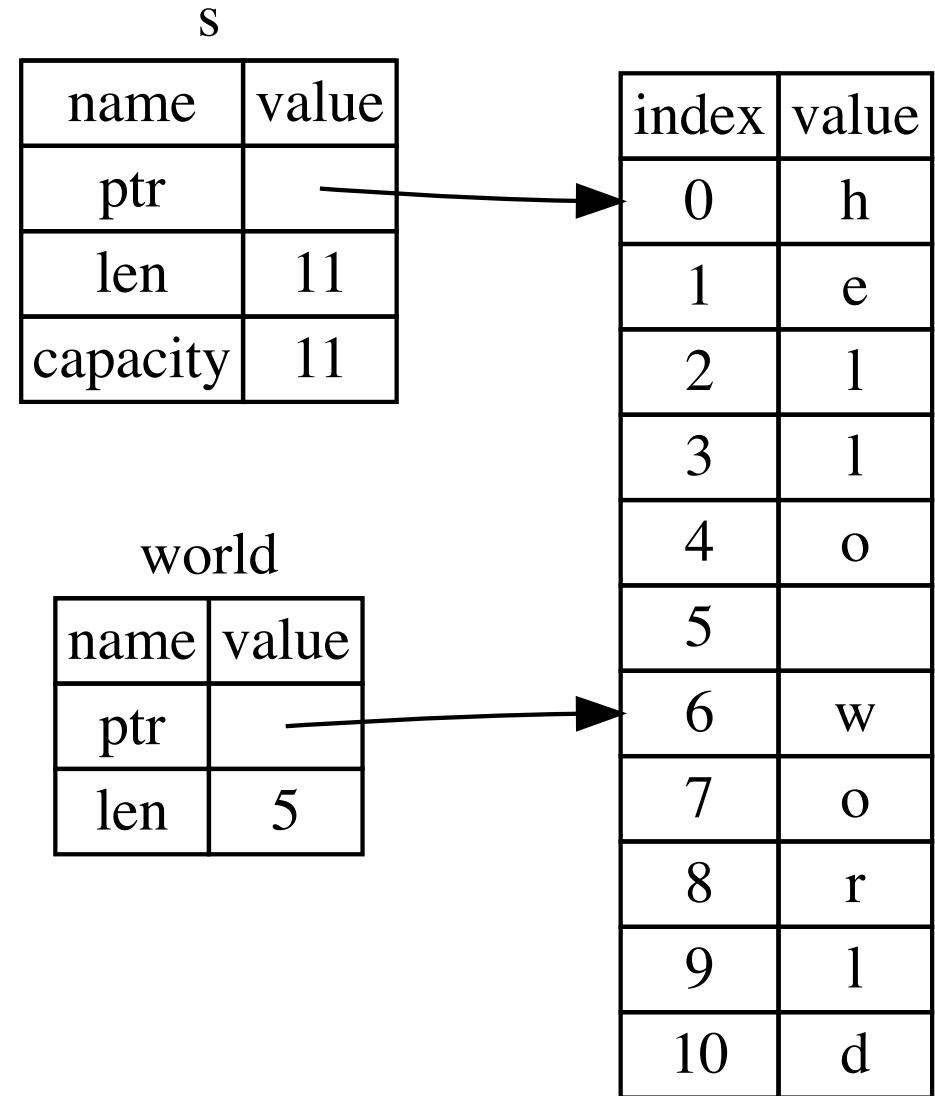
```
let hello = &s[0..5];
let world = &s[6..11];
```

- `hello` contains the first 5 characters of `s`
- `world` contains the 5 characters starting at the 6th index of `s`

String Slices

```
let s = String::from("hello world");  
  
let hello = &s[0..5];  
let world = &s[6..11];
```

- A string slice stores a pointer to memory and a length



String Slices

You can shorthand ranges with the `...` syntax.

```
let s = String::from("hello");

let slice = &s[0..2];
let slice = &s[..2];

let len = s.len();
let slice = &s[3..len];
let slice = &s[3..];

let slice = &s[0..len];
let slice = &s[...];
```



String Literals are Slices

Recall that we talked about string literals being stored inside the binary.

```
let s: &str = "Hello, world!";
```

- The type of `s` here is `&str`: it's a slice pointing to that specific point of the binary with type `str`
- String literals are immutable
 - Their `&str` immutable reference type reflects that

Owned Types

- String slices and string literals are immutable because they are a special type of immutable reference
- String is an owned type
 - i.e. a type that has an owner
- Another owned type is a *vector*

Vectors

A *vector* allows you to store a collection of values contiguously in memory.

You can create a vector with the method `new`:

```
let v: Vec<i32> = Vec::new();
```

- Internally, a `Vec` is a dynamically sized array stored on the heap
- All values in the `Vec` must be of the same type
- The `<i32>` just means that the vector stores `i32` values
 - We'll talk more about this `<>` syntax in week 4!

Updating a `Vec`

To add elements to a `Vec`, we can use the `push` method.

```
let mut v = Vec::new();  
  
v.push(5);  
v.push(6);  
v.push(7);  
v.push(8);  
  
println!("{:?}", v);
```

```
[5, 6, 7, 8]
```

vec! Macro

Rust provides a *macro* to create vectors easily in your programs.

```
let v = vec![1, 2, 3];  
println!("{:?}", v);
```

```
[1, 2, 3]
```

- Briefly: Macros are a special type of function
 - They can take in a variable number of arguments

Reading Elements of Vectors

You can index into a vector to retrieve a reference to an element.

```
let v = vec![1, 2, 3, 4, 5];  
  
let third: &i32 = &v[2];  
println!("The third element is {}", third);
```

- Note that Rust will panic if you try to index out of the bounds of the `Vec`

More `Vec<T>` to come...

We will talk more about `String` and `Vec<T>` in week 4!

Homework 2

- The second homework consists of 12 small ownership puzzles
 - Refer to the `README.md` for further instructions
 - Always follow the compiler's advice!
- We *highly* recommend reading the Rust Book chapter on [ownership](#)
 - Ownership is a *very tricky concept* that affects almost every aspect of Rust, so understanding it is key to writing more complex Rust code
- Try your best to understand Ownership *before* attempting the homework

Next Lecture: Structs and Enums

Thanks for coming!

Slides created by:

Connor Tsui, Benjamin Owad, David Rudo,
Jessica Ruan, Fiona Fisher, Terrance Chen

