



# *Intro to Rust Lang*

# **Unsafe**

# The Story So Far...

- We have covered all of the basic features of Rust, as well as many of the intermediate concepts
- If you are confident you understand the past 10 lectures, you can say you are proficient with Rust!

# Epilogue

As much as we'd have loved to dive deep into all of these topics, we simply do not have that much time.

However...

- The goal of this course was never to feed you information
- The goal was to teach you the *core ideas* of Rust and how to think about it
- We hope that you will take the knowledge from this class and use it to explore more about this programming language *yourself*

# Unsafe Rust

# Into the Woods

So far, we've only seen code where memory safety is guaranteed at compile time.

- Rust has a second language hidden inside called *unsafe Rust*
- `unsafe` Rust does not enforce memory safety guarantees

# Why unsafe?

- Static analysis is *conservative*
- By definition, it enforces *soundness* rather than *completeness*
- We need a way to tell the compiler: "Trust me, I know what I'm doing"
- Additionally, computer hardware is inherently unsafe
  - *Network goes down...*
  - *Someone unplugs your hard drive...*
  - *Bit flip from cosmic ray...*

## unsafe in 2025

- Rust's precise requirements for `unsafe` code are still being determined
- There's an entire book dedicated to `unsafe` Rust called the [Rustonomicon](#)

## What is **unsafe**, really?

If you take anything away from today, it should be this:

**Unsafe code is the mechanism Rust gives developers for taking advantage of invariants that, for whatever reason, the compiler cannot check.**

- *Jon Gjengset, Rust for Rustaceans*

## What `unsafe` is not

It's important to understand that `unsafe` is *not* just a way to skirt the rules of Rust. Recall the rules of Rust:

- Ownership
- Borrow Checking
- Lifetimes
- `unsafe` is a way to *enforce* these rules using reasoning beyond the compiler
  - The onus is on *you* to ensure the code is **safe**
  - *This is usually the case with memory-unsafe languages!*

## Other Languages

You will likely not have to deal with `unsafe` unless you are dealing with low-level systems primitives and concepts.

However, a lot of the concepts you will learn in this lecture can be applied to other languages, even if people don't usually explicitly say so.

- Understanding undefined behavior is *super* important in C and C++

# The `unsafe` Keyword

There are 2 ways to use the `unsafe` keyword in Rust. The first is marking a *function* as `unsafe`.

```
impl<T> SomeType<T> {
    // vvvvvv
    pub unsafe fn decr(&self) {
        self.some_usize -= 1;
    }
}
```

- Here, the `unsafe` keyword serves as a warning to the caller
- Notice how the function body is totally normal!
- There may be additional invariants that must be upheld *before* calling `decr`

## The `unsafe` Keyword

The second way is marking an *expression* as `unsafe`.

```
impl<T> SomeType<T> {
    pub fn as_ref(&self) -> &T {
        unsafe { &*self.ptr }
    }
}
```

# The `unsafe` Contracts

```
impl<T> SomeType<T> {
    pub unsafe fn decr(&self) {
        self.some_usize -= 1;
    }

    pub fn as_ref(&self) -> &T {
        unsafe { &*self.ptr }
    }
}
```

- The first requires the caller to be careful
- The second assumes the caller *was* careful when invoking `decr`

# The unsafe Contracts

Imagine if `SomeType<T>` was really `Rc<T>`:

```
impl<T> Rc<T> {
    pub unsafe fn decr_ref_count(&self) {
        self.count -= 1;
    }

    pub fn as_ref(&self) -> &T {
        unsafe { &*self.ptr }
    }
}
```

- When `self.count` hits 0, `T` is dropped
- What if someone else constructed `&T` without incrementing `self.count`?
- As long as nobody corrupts the reference count, this code is safe

# Unsafe Superpowers

So what can we do with `unsafe` ?

With `unsafe` , we get 5 superpowers! We can:

1. Call an `unsafe` function or method
2. Access or modify a mutable static variable
3. Implement an `unsafe` trait
4. Access fields of `union`s

# Unsafe Superpowers

1. Call an `unsafe` function or method
2. Access or modify a mutable static variable
3. Implement an `unsafe` trait
4. Access fields of `union`s

These 4 things aren't all that interesting, so why the big fuss?

# THE UNSAFE SUPERPOWER

The **biggest** superpower of all is superpower 5!

- DEREference A RAW POINTER
  - That's it!
  - *But honestly, it's enough to wreak all sorts of havoc...*

# Raw Pointers

Unsafe Rust has 2 types of Raw Pointers:

- `*const T` is an immutable raw pointer
- `*mut T` is a mutable raw pointer
- *Note that the asterisk `*` is part of the type name*
- *Immutable* here means that the pointer can't be reassigned directly after being dereferenced

# Pointers vs References

Raw Pointers themselves are allowed to do some special things:

- They can ignore borrowing rules by having multiple immutable and mutable pointers to the same location
- They are not guaranteed to point to valid memory
- They don't implement any automatic cleanup
- They can be `NULL` 

# Raw Pointers Example

Here's an example of creating raw pointers:

```
let mut num = 5;  
  
let r1 = &num as *const i32;  
let r2 = &mut num as *mut i32;
```

- We have both an immutable and mutable pointer pointing to the same place
- Notice how there is no `unsafe` keyword here
- We can *create* raw pointers safely, we just cannot *dereference* them

# Raw Pointers Example

Here is another example of creating a raw pointer:

```
let address: usize = 0xDEADBEEF;  
let r = address as *const i32;
```

- We construct a pointer to (likely invalid) memory
- Again, no `unsafe` keyword necessary here!

# Raw Pointers and `unsafe`

Let's actually try and dereference these pointers:

```
let mut num = 5;

let r1 = &num as *const i32;
let r2 = &mut num as *mut i32;

unsafe {
    println!("r1 is: {}", *r1);
    println!("r2 is: {}", *r2);
}
```

- There's no undefined behavior here? Right?
- *Right?*
- Right! 🦀



# Calling `unsafe` Functions

We must also call `unsafe` functions in an `unsafe` block.

```
unsafe fn dangerous() {}

fn main() {
    unsafe {
        dangerous();
    }
}
```

- We would get an error if we called `dangerous` without the `unsafe` block!

# Using `extern` Functions

Sometimes, we might need to interact with code from another language.

- Rust has the keyword `extern` that facilitates the use of a *Foreign Function Interface (FFI)*
- Since other languages do not have Rust's safety guarantees, we have no idea if they are safe to call or not!

## extern "C"

Let's see how we would set up integration with the `abs` function from the C standard library.

```
extern "C" {
    fn abs(input: i32) -> i32;
}

fn main() {
    unsafe {
        println!("Absolute value of -3 according to C: {}", abs(-3));
    }
}
```

# extern "C"

```
extern "C" {
    fn abs(input: i32) -> i32;
}

fn main() {
    unsafe {
        println!("Absolute value of -3 according to C: {}", abs(-3));
    }
}
```

- The `extern "C"` declares the *Application Binary Interface (ABI)* that this foreign function uses
- We have no idea if `abs` is doing what it is supposed to be doing
  - It is on us as the programmer to ensure safety



## extern "C"

We can also use `extern` to allow other languages to call Rust code!

```
#[no_mangle]
pub extern "C" fn call_from_c() {
    println!("Just called a Rust function from C!");
}
```

- Note that this is *not* a foreign function, this us "exporting" our Rust function to other languages
- Also observe how the usage of `extern` does not require `unsafe`

# Mutable Static Variables

We can mutate global static variables with `unsafe`.

```
static mut COUNTER: u32 = 0;

fn add_to_count(inc: u32) {
    unsafe {
        COUNTER += inc;
    }
}

fn main() {
    add_to_count(3);
    unsafe {
        println!("COUNTER: {}", COUNTER);
    }
}
```

## Last 2 Superpowers

The last 2 superpowers are implementing an `unsafe` trait and accessing fields of a `union`.

- `Send` and `Sync` are both `unsafe` traits
  - The developer must provide their own proof of thread safety
  - *More on thread safety in the next lecture!*
- `union`s are primarily used to interface with unions in C code

## How to use `unsafe` code

- Just because a function contains `unsafe` code doesn't mean you need to mark the entire function as `unsafe`
- Often, we want to write `unsafe` code that we *know* is actually safe
- A common abstraction is to wrap `unsafe` code in a safe function

## split\_at\_mut

Let's take a look at `split_at_mut` from the standard library.

```
let mut v = vec![1, 2, 3, 4, 5, 6];  
  
let r = &mut v[..];  
  
let (a, b) = r.split_at_mut(3);  
  
assert_eq!(a, &mut [1, 2, 3]);  
assert_eq!(b, &mut [4, 5, 6]);
```

## split\_at\_mut

```
fn split_at_mut(values: &mut [i32], mid: usize) -> (&mut [i32], &mut [i32]);
```

- Unfortunately, we cannot write this function using only safe Rust
- How would we attempt it?

# split\_at\_mut Implementation

Here is a trivial first attempt:

```
fn split_at_mut(values: &mut [i32], mid: usize) -> (&mut [i32], &mut [i32]) {  
    let len = values.len();  
  
    assert!(mid <= len);  
  
    (&mut values[..mid], &mut values[mid..])  
}
```

- What is the issue with this?
- Can you figure out what the compiler will tell us *just by looking at the function signature?*



# split\_at\_mut Compiler Error

If we try to compile, we get this error:

```
error[E0499]: cannot borrow `*values` as mutable more than once at a time
--> src/main.rs:6:31
1 | fn split_at_mut(values: &mut [i32], mid: usize) -> (&mut [i32], &mut [i32]) {
|   - let's call the lifetime of this reference ``1``
...
6 |     (&mut values[..mid], &mut values[mid..])
|-----^~~~~~-----|
|   |           |
|   |           second mutable borrow occurs here
|   |           first mutable borrow occurs here
|   returning this value requires that `*values` is borrowed for ``1``
```

For more information about this error, try `rustc --explain E0499`.  
error: could not compile `unsafe-example` due to previous error



# split\_at\_mut Implementation

Let's try again with `unsafe`.

```
fn split_at_mut(values: &mut [i32], mid: usize) -> (&mut [i32], &mut [i32]) {
    let len = values.len();
    let ptr = values.as_mut_ptr();

    assert!(mid <= len);

    unsafe {
        (
            std::slice::slice::from_raw_parts_mut(ptr, mid),
            std::slice::slice::from_raw_parts_mut(ptr.add(mid), len - mid),
        )
    }
}
```

## split\_at\_mut Implementation

`from_raw_parts_mut` is `unsafe` because it takes a raw pointer and length, and it must trust that these form a valid slice.

```
unsafe {
(
    std::slice::slice::from_raw_parts_mut(ptr, mid),
    std::slice::slice::from_raw_parts_mut(ptr.add(mid), len - mid),
)
}
```

- Since the `ptr` came from a valid slice, we know the first is valid
- Since we know `[mid, len)` is valid memory, we also know the second is valid!

# from\_raw\_parts\_mut Safety Contract

Here is the actual safety contract for `from_raw_parts_mut` :

```
/// # Safety
///
/// Behavior is undefined if any of the following conditions are violated:
///
/// * `data` must be [valid] for both reads and writes for `len * mem::size_of::<T>()` many bytes,
///   and it must be properly aligned. This means in particular:
///
///     * The entire memory range of this slice must be contained within a single allocated object!
///       Slices can never span across multiple allocated objects.
///     * `data` must be non-null and aligned even for zero-length slices. One
///       reason for this is that enum layout optimizations may rely on references
///       (including slices of any length) being aligned and non-null to distinguish
///       them from other data. You can obtain a pointer that is usable as `data`
///       for zero-length slices using [`NonNull::dangling()`].
///
/// * `data` must point to `len` consecutive properly initialized values of type `T`.
///
/// * The memory referenced by the returned slice must not be accessed through any other pointer
///   (not derived from the return value) for the duration of lifetime `a`.
///   Both read and write accesses are forbidden.
///
/// * The total size `len * mem::size_of::<T>()` of the slice must be no larger than `isize::MAX`.
```



## from\_raw\_parts\_mut Misuse

We could very easily misuse `from_raw_parts_mut` if we wanted to...

```
use std::slice;

let address: usize = 0xDEADBEEF;
let r = address as *mut i32;

let values: &[i32] = unsafe { slice::from_raw_parts_mut(r, 1000) };
```

- This might seem ridiculous, but when you always assume your code is safe...

# SAFETY Contracts

Whenever you use `unsafe`, make sure to write proof of why your code is not undefined behavior.

```
let mut v: Vec<i32> = vec![1, 2, 3, 4];
// SAFETY: We know that `v` has 4 elements, so adding `3 * size_of::<i32>()` ` 
// cannot overflow (wrap around the address space).
let p: *const i32 = unsafe { v.as_ptr().add(3) };

v.push(5);

// SAFETY: We have checked that a resize will never happen with only
// 5 elements, so this dereference is safe.
let elem: i32 = unsafe { *p };
println!("{}", elem);
```

- This is simplified!

Note that this is what the `unsafe` code should *really* look like:

```
let mut v: Vec<i32> = vec![1, 2, 3, 4];

// SAFETY: We know that `v` has 4 valid elements, which means that the pointers at
// offsets +1, +2, and +3 are all valid pointers that point to values of type `i32`.
// This also means it cannot be the case that adding `3 * size_of::<i32>()` to
// `v.as_ptr()` overflows, otherwise that last element could not even exist in the
// first place. With these two guarantees, we fulfill `add`'s safety contract.
let p: *const i32 = unsafe { v.as_ptr().add(3) };

v.push(5);

// SAFETY: We showed above that `p` was a valid pointer to an `i32` value. We have
// also checked in the code of `Vec::push` that a resize and reallocation
// **cannot** happen when there are only 5 elements. So since there is no
// reallocation, and `v` is still in scope (so the memory has not been freed), this
// dereference of `p` to an `i32` value is valid.
let elem: i32 = unsafe { *p };

println!("{}", elem);
```



# With Great Power...

What could go wrong?

- Probably not much, *if* you're careful
  - By careful, we mean writing a proof for every use of `unsafe`
- If you do get something wrong...
- With `unsafe`, you hold great responsibility

# Undefined Behavior

If you get something wrong, your program now has *undefined behavior*.

- It should go without saying that undefined behavior is bad...
- The best scenario is you get a visible error:
  - Segfaults
  - Unexpected deadlocks
  - Garbled output
  - Panics that *don't* exit the program
- The worst case...

# Undefined Behavior

The worst case scenario is that your program state is *invisibly* corrupted.

- Data races
- Transactions aren't atomic
- Backups are corrupted
- Security leaks
- Schrödinger's Bug

# Interacting with Safe Rust

Unsafe code is not defined.

- The compiler could eliminate the entire `unsafe` block if it wanted to
- It could also miscompile surrounding, safe code!
- In a lot of ways, `unsafe` Rust is far worse than C/C++ because it assumes *all* of Rust's safety guarantees

## Safe **unsafe**: Valid References

You may recall that all references must be valid. A valid reference:

- Must never dangle
- Must always be aligned
- Must always point to a valid value for their target type
- Must either be immutably shared or mutably exclusive
- Plus more guarantees relating to lifetimes
- If you create a reference from a pointer, make sure all of these are true!



# Other Validity Requirements

Some primitive types have other guarantees:

- `bool` is 1 byte, but can only hold `0x00` or `0x01`
- `char` cannot hold a value above `char::MAX`
- Most Rust types cannot be constructed from uninitialized memory
- If Rust didn't enforce this, it wouldn't be able to make niche optimizations
  - `Option<&T>` is a good example
  - What if `Option<Option<bool>>` used `0x00` through `0x03` ?
- It doesn't really matter if Rust does or does not make the optimization
  - All that matters is that it is *allowed* to whenever it wants

# Even More Validity Requirements

Here are some even more requirements:

- Owned Pointer Types (like `Box` and `Vec`) are subject to optimizations assuming the pointer to memory is not shared or aliased anywhere
- You can never assume the layout of a type when casting
- All code must be prepared to handle `panic!`s and *stack unwinding*
- Stack unwinding drops everything in the current scope, returns from that scope, drops everything in that scope, returns, etc...
- All variables are subject to something called the *Drop Check*, and if you drop something incorrectly, you might cause undefined behavior

# Fighting with `unsafe`

That was a lot, right?

- Remember that it is very possible to write safe `unsafe` code
- A lot of the time, it isn't actually that difficult
- Being careful is half the battle
- Being absolutely sure you actually need `unsafe` is the other half

# Working with `unsafe`

It is tempting to reason about unsafety *locally*.

- Consider whether the code in the `unsafe` block is safe in the context of both the rest of the codebase, and in the context of other people using your library
- Encapsulate the unsafety as best you can
- Read and write documentation!
- Use tools like `Miri` to verify your code!
- **Make sure to formally reason about your program**

# Miri

[Miri](#) is an undefined behavior detection tool for Rust.

- An interpreter for Rust's mid-level intermediate representation
- Can detect out-of-bounds memory accesses and use-after-free
- Invalid use of uninitialized data
- Not sufficiently aligned memory accesses and references
- Can also detect data races
- Think of Miri as Valgrind, address sanitizer, and thread sanitizer, all in one!

## Recap: `unsafe`

- With `unsafe`, we have great powers
- But we must accept the responsibility of leveraging those powers
- There are consequences to writing unsafe `unsafe` code
- `unsafe` is a way to *promise* to the compiler that the indicated code is safe

## Next Lecture: Parallelism

Thanks for coming!

*Slides created by:*

Connor Tsui, Benjamin Owad, David Rudo,  
Jessica Ruan, Fiona Fisher, Terrance Chen

