



***Intro to Rust Lang***  
***Standard***  
***Collections and***  
***Generics***

# Today: Standard Collections and Generics

- Rust's `std::collection` Types
  - `Vec<T>`
  - `String`
  - `HashMap<K, V>`
- Generic Types

# Standard Collections

Rust's standard library contains a number of useful data structures called *collections*.

- Most other data types represent a single value, but collections contain multiple values
- Values in collections are (*almost always*) stored on the **heap**
  - The size of each collection does not need to be known at compile time
  - *For more information on other standard library collections, refer to the [documentation](#) of the `std::collections` module*

# Vectors

# Review: Vectors

You can create a vector with `new`, and add elements with `push`.

```
let mut v = Vec::new();  
  
v.push(5);  
v.push(6);  
v.push(7);  
v.push(8);  
  
println!("{:?}", v);
```

## Review: `vec!` Macro

Rust provides a *macro* to create vectors easily in your programs.

```
let v = vec![1, 2, 3];  
  
println!("{:?}", v);
```

```
[1, 2, 3]
```

## Review: Indexing

You can index into a vector to retrieve a reference to an element.

```
let v = vec![1, 2, 3, 4, 5];

let third_ref: &i32 = &v[2];
println!("The third element is {}", third_ref);

let third: i32 = v[2]; // This is only possible because i32 is Copy
println!("The third element is {}", third);
```

## Vec::get()

You can also use the `get` method to access an optional reference.

```
let v = vec![1, 2, 3, 4, 5];

let third: Option<&i32> = v.get(2);
match third {
    Some(third) => println!("The third element is {}", third),
    None => println!("There is no third element."),
}
```

- Using `get` returns `None` if the index is out of bounds instead of panicking



# Vec and References

Recall the rules for immutable and mutable references.

```
let mut v = vec![1, 2, 3, 4, 5];  
  
let first = &v[0];  
  
v.push(6);  
  
println!("The first element is: {}", first);
```

- You cannot mutate a vector while references to its elements exist
- Appending might resize and reallocate the vector and change its location in memory



# Vec and References

If we try to run this:

```
let mut v = vec![1, 2, 3, 4, 5];
let first = &v[0];
v.push(6);
println!("The first element is: {}", first);
```

```
error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable
--> src/main.rs:4:5
   |
3  |     let first = &v[0];
   |                   - immutable borrow occurs here
4  |     v.push(6);
   |     ^^^^^^^^^^^ mutable borrow occurs here
5  |     println!("The first element is: {}", first);
   |                                           ----- immutable borrow later used here
```

# Iterating over a Vector

To access each element in order, we can iterate through the elements with a `for` loop directly, rather than using indices.

```
let v = vec![100, 32, 57];

for elem in &v { // `elem` is a reference to an i32 (&i32)
    println!("{}", elem);
}
```

```
100
32
57
```

# Mutable iteration over a Vector

We can also iterate over mutable references to each element to make changes to each element.

```
let mut v = vec![100, 32, 57];

for elem in &mut v { // `elem` is a mutable reference to an i32
    *elem += 50;
}

println!("{:?}", v);
```

```
[150, 82, 107]
```

- We only have a single mutable reference into the vector at a time
  - We pass the borrow checker's rules!

# For Loop Sugar

You can also *consume* the vector when you want to loop over it.

```
let v = vec![100, 32, 57];

for i in v {
    println!("{}", i);
}

// println!("{:?}", v); <-- Can't do this anymore!
```

- We'll talk more about this in week 7!

## Deref Coercion to `&[T]`

Instead of a function taking a `&Vec<T>` as a parameter, we can take a `&[T]`.

```
fn largest(list: &Vec<i32>) -> &i32
```

```
fn largest(list: &[i32]) -> &i32
```

- The second is more general and preferred
- We can do this because of *deref coercion*
  - This basically means you can turn a `&Vec<T>` into a `&[T]`
- We'll talk more about this in week 9!

# Use Enums to Store Multiple Types

Vectors can only store values of the same type, so we can use enums to store values of different types (variants).

```
enum SpreadsheetCell {
    Int(i32),
    Float(f64),
    Text(String),
}

let row = vec![
    SpreadsheetCell::Int(3),
    SpreadsheetCell::Text(String::from("blue")),
    SpreadsheetCell::Float(10.12),
];
```

# Vectors and Ownership

Vectors own all of their contained elements.

```
let mut v = vec![String::from("rust"), String::from("is")];  
  
let s = String::from("great!");  
  
v.push(s); // move `s` into `v`  
  
// `s` is no longer usable here!
```

- To insert an owned value, it must be *moved* into the vector
  - In other words, ownership must be transferred to the vector



# Dropping a Vector

Like any other struct, a vector is dropped when it goes out of scope.

```
{  
    let v = vec!["rust".to_string(), "is".to_string(), "great!".to_string()];  
    stuff(&v); // do stuff with `v`  
} // <- `v` goes out of scope and everything it contains is freed
```

- When the vector gets dropped, all of its contents are also dropped
- The borrow checker will ensure that references into the vector cannot be used after it has been dropped

**String**

# What is a String?

- A `String` is essentially a `Vec` of bytes interpreted as text
- We introduced them back in week 2, but now we'll look at them in more depth
- New Rust programmers may be confused by:
  - `String`'s internal UTF-8 encoding
  - Rust's prevention of possible logical errors from the encoding
  - `String`s are not as simple as they may initially seem

# What is a String?

- Rust "only" has one string type in the core language, `str`
  - We almost always see it in its borrowed form, `&str`
  - String slices are `&str`
  - String literals are also `&str`
    - References data stored in the program's binary
- `String` is a growable, mutable, owned, UTF-8 encoded string type

# Creating a `String`

You can create a `String` with the methods `new`, `to_string`, or `from`.

```
let mut s = String::new(); // empty mutable string

let data = "initial contents"; // string literal

let s = data.to_string(); // string literal into `String`

// the method also works on a literal directly:
let s = "initial contents".to_string();

let s = String::from("initial contents"); // string literal into `String`
```

# Strings are UTF-8 Encoded

We can represent any properly encoded data in `String`.

Here are some greetings in different languages!

```
let hello = String::from("السلام سييم");  
let hello = String::from("Dobrý den");  
let hello = String::from("Hallo");  
let hello = String::from("שלום");  
let hello = String::from("नमस्ते");  
let hello = String::from("Olá");  
let hello = String::from("Привет");  
let hello = String::from("Hola");
```

# Updating a `String`

We can grow a `String` by using the `push` or `push_str` methods.

```
let mut s = String::from("foo");  
  
s.push('b'); // push a char  
s.push_str("ar"); // push a &str  
  
println!("{}", s);
```

```
foobar
```

## Updating a `String`

The `push_str` method takes a string slice, because we don't want to take ownership of the string passed in.

```
let mut s1 = String::from("foo");  
  
let s2 = String::from("bar");  
  
s1.push_str(&s2);  
  
println!("s2 is {}", s2); // `s2` is still valid!
```

```
s2 is bar
```



# Concatenating Strings

You can combine two strings with `+`:

```
let s1 = String::from("Hello, ");
let s2 = String::from("world!");
let s3 = s1 + &s2; // note s1 has been moved here and can no longer be used
```

This is syntactic sugar for a function whose signature looks something like this:

```
fn add(self, s: &str) -> String
```

- Notice that `add` takes full ownership of `self`
- Also notice `add` takes `&str` as its second parameter and not `&String`
  - This is the same *deref coercion* as with `&Vec<T>` to `&[T]`!

# Concatenating Multiple Strings

```
let s1 = String::from("tic");  
let s2 = String::from("tac");  
let s3 = String::from("toe");
```

To combine multiple strings, you can use multiple `+`'s:

```
let s = s1 + "-" + &s2 + "-" + &s3;
```

Or you can use the `format!` macro:

```
let s = format!("{}", s1, s2, s3);  
let s = format!("{s1}-{s2}-{s3}"); // relatively new shorthand!
```

# Indexing into Strings

This code might seem reasonable from any other programming language like Python or C.

```
let s1 = String::from("hello");  
let h = s1[0];
```

- Accessing individual characters in a string by indexing is common in many languages
- However, if you try to access a `String` using an index, you will get an error



# Indexing into Strings

```
let s1 = String::from("hello");  
let h = s1[0];
```

```
error[E0277]: the type `String` cannot be indexed by `{integer}`  
--> src/main.rs:3:13  
3 |     let h = s1[0];  
  |               ^^^^^ `String` cannot be indexed by `{integer}`  
= help: the trait `Index<integer>` is not implemented for `String`
```

- Why won't Rust allow indexing `String`?

# Internal Representation of Strings

A `String` is really a wrapper over `Vec<u8>`, or a vector of bytes.

```
let hello = String::from("Hola");
```

- How long is this string?
  - The length of the string is 4
  - The internal vector storing the string `"Hola"` is 4 bytes long
- Simple enough, right?

# Internal Representation: Cyrillic

Now suppose we wanted to say "Hello", but in Russian.

```
let hello = String::from("Привет");
```

- How long is this string?
  - There are 6 distinct characters
  - However, the string's `len` is 12, the number of bytes needed in the internal vector

# Internal Representation: UTF-8

Let's revisit some invalid Rust code again.

```
let hello = "Привет";  
let answer = &hello[0];
```



- What *should* `answer` be?
  - It can't be `п`, internally it is represented by 2 bytes: `[208, 159]`
  - Do we return `208` instead?
- There isn't any obvious expected behavior here...

# Internal Representation: UTF-8

```
let hello = "Привет";  
let answer = &hello[0]; // BAD!
```

- Anything we can return here might not be an "expected" result
- The philosophy of Rust is to not compile this code at all
  - Prevents misunderstandings early in the development process
- Further reading on UTF-8: [Rust Book Chapter 8.2](#)



# Slicing Strings

Instead of indexing with a single number, you can use `[]` with a *range* to create a string slice containing specific bytes.

```
let hello = "Привет";  
let s = &hello[0..4]; // `s` == "Пр"
```

# Slicing Strings

However, if we try to slice only a part of a character's bytes, Rust panics at runtime.

```
let hello = "Привет";  
  
let s = &hello[0..1];
```

```
thread 'main' panicked at 'byte index 1 is not a char boundary;  
it is inside 'П' (bytes 0..2) of `Привет`'
```

- This happens in the same way that an invalid index causes a panic!

# Iterating Over Strings

Normally, we want to iterate over individual Unicode scalar values, and we can use the `chars` method.

```
for c in "Πp".chars() {  
    println!("{c}");  
}
```

```
Π  
p
```

# Iterating Over Strings

Alternatively, if you want the actual raw bytes, you can use the `bytes` method.

```
for b in "Пр".bytes() {  
    println!("{b}");  
}
```

```
208  
159  
209  
128
```

## Recap: Strings

- Rust chooses to use UTF-8 strings as the default (for both `String` and `&str`)
  - Programmers have to think about handling unicode upfront
  - The complexity brought on by encodings is more apparent in Rust
  - However, this prevents having to deal with non-ASCII characters later!
- The standard library offers many methods for `String` and `&str` types to help handle these complex situations correctly

# HashMap

# HashMap<K, V>

The type `HashMap<K, V>` stores keys with type `K` mapped to values with type `V`.

- Many languages support this kind of data structure, even if they use a different name:
  - Hash
  - Map
  - Object
  - Hash Table
  - Dictionary
  - Associative Array

# Creating a Hash Map

We can create a new hash map with `new` and insert entries with `insert`.

```
use std::collections::HashMap;

let mut scores = HashMap::new();

scores.insert(String::from("Blue"), 10);
scores.insert(String::from("Yellow"), 50);
```

- Note that we need to import `HashMap` from the standard library's `collections` module with `use`
- We'll talk more about `use` in week 6!



# Accessing Values in a Hash Map

We can use the `get` method to get the value associated with a key.

```
let mut scores = HashMap::new();
scores.insert(String::from("Blue"), 10);
scores.insert(String::from("Yellow"), 50);

let team_name = String::from("Blue");
let score = scores.get(&team_name).unwrap_or(&0);
```

- The `get` method returns an `Option<&V>`, similar to `Vec::get()`
- We use `unwrap_or(&0)` on the result
  - If it returns `Some(&x)`, we unwrap and get `&x`
  - If it returns `None`, we go to the default case `&0`

# Iterating over a Hash Map

We can iterate over each key/value pair using a `for` loop, similar to vectors.

```
let mut scores = HashMap::new();

scores.insert(String::from("Blue"), 10);
scores.insert(String::from("Yellow"), 50);

for (key, value) in scores {
    println!("{key}: {value}");
}
```

```
Yellow: 50
Blue: 10
```

- *Note that the order is non-deterministic*

# Hash Maps and Ownership

Hash maps own their contained data, just like vectors.

```
let field_name = String::from("Favorite color");  
let field_value = String::from("Blue");  
  
let mut map = HashMap::new();  
map.insert(field_name, field_value);  
  
// field_name and field_value are invalid at this point!
```

# Updating a Hash Map

Hash maps only contain one value for each distinct key, so to update we can just insert twice.

```
let mut scores = HashMap::new();  
  
scores.insert(String::from("Blue"), 10);  
scores.insert(String::from("Blue"), 25);  
  
println!("{:?}", scores);
```

```
{"Blue": 25}
```

- Inserting twice overwrites the existing value for the given key

# Accessing a Hash Map with Defaults

- A common pattern when accessing a `HashMap` is:
  - If the key exists, we want to access the value
  - If the key does not exist, insert a default value and then access it
- `HashMap` has a special API called `Entry`

# Hash Map Entries

To insert a value if the key does not already exist, you can use the `Entry` API and the method `or_insert`.

```
let mut scores = HashMap::new();
scores.insert(String::from("Blue"), 10);

scores.entry(String::from("Yellow")).or_insert(50);
scores.entry(String::from("Blue")).or_insert(50);

println!("{:?}", scores);
```

```
{"Yellow": 50, "Blue": 10}
```

# Hash Map Entries

If you want to update a value, or provide a default if it doesn't yet exist, you can do something similar:

```
let text = "hello world wonderful world";

let mut map = HashMap::new();

for word in text.split_whitespace() {
    let count: &mut usize = map.entry(word).or_insert(0);
    *count += 1;
}

println!("{:?}", map);
```

```
{"world": 2, "hello": 1, "wonderful": 1}
```

# hash\_map::Entry::or\_insert

The method `or_insert` has the following signature:

```
fn or_insert(self, default: V) -> &mut V
```

- It gives out a mutable reference
  - That reference are guaranteed to point to valid data
  - We need to provide a default, otherwise this data might not exist
- Shorter and more readable than separate conditionals



# Recap:

- We covered [The Rust Book Chapter 8](#)
- Always refer to the documentation!
  - [Vec<T> documentation](#)
  - [String documentation](#)
  - [HashMap<K, V> documentation](#)

# Generics

# Generics

So what was the deal with the `T` in `Vec<T>`, and `K, V` in `HashMap<K, V>`?

- We refer to these as *generic* types
- Think of it as being able to fill in any type you want in place of `T`
- Generics allow us to replace specific types with a placeholder that represents multiple types
  - Removes code duplication

# Removing Code Duplication

Let's say we want to find the largest number in a list.

```
let number_list = vec![34, 50, 25, 100, 65];

let mut largest = &number_list[0];

// Pretend the list always has at least 1 element.
for number in &number_list {
    if number > largest {
        largest = number;
    }
}

println!("The largest number is {}", largest);
```

# Removing Code Duplication

What if we have multiple lists? We then have to do multiple searches.

```
let number_list = vec![34, 50, 25, 100, 65];
let mut largest = &number_list[0];
for number in &number_list {
    if number > largest {
        largest = number;
    }
}
println!("The largest number is {}", largest);

let number_list = vec![102, 34, 6000, 89, 54, 2, 43, 8];
let mut largest = &number_list[0];
for number in &number_list {
    if number > largest {
        largest = number;
    }
}
```

# Removing Code Duplication

Instead, we can make a function called `largest`.

```
fn largest(list: &[i32]) -> &i32 {
    let mut largest = &list[0];
    for item in list {
        if item > largest {
            largest = item;
        }
    }
    largest
}

fn main() {
    let number_list = vec![34, 50, 25, 100, 65];
    println!("The largest number is {}", largest(&number_list));

    let number_list = vec![102, 34, 6000, 89, 54, 2, 43, 8];
    println!("The largest number is {}". largest(&number_list));
}
```

# Remove Function Duplication

What if we also wanted to find the largest character in a slice?

```
fn largest_char(list: &[char]) -> &char {
    let mut largest = &list[0];
    for item in list {
        if item > largest {
            largest = item;
        }
    }
    largest
}
```

- This is effectively the same as finding the largest number in a list
- We would still need to write a new function in addition to `largest`
- Can we remove a *function* that has been duplicated?

# Generic Functions

We can define a function as generic over some type `T` with a tag `<T>`:

```
fn largest<T>(list: &[T]) -> &T
```

- This function is generic over `T`
- This function takes in a slice of `T` as input
- This function returns a reference to `T`
- `T` can be *any*\* type!



# Generic Functions

Generic types can have any name, not just `<T>`:

```
fn largest<T>(list: &[T]) -> &T
```

```
fn largest<Key>(list: &[Key]) -> &Key
```

```
fn largest<Smile>(list: &[Smile]) -> &Smile
```

- All of these essentially mean the same thing!
  - Last one is *frowned* upon since it might seem like a struct or enum

# Generic Functions

Let's try to modify our old function directly:

```
fn largest<T>(list: &[T]) -> &T {
    let mut largest = &list[0];
    for item in list {
        if item > largest {
            largest = item;
        }
    }
    largest
}

fn main() {
    println!("The largest number is {}",
        largest(&[34, 50, 25, 100, 65]));
    println!("The largest char is {}",
        largest(&['y', 'm', 'a', 'q']));
}
```



# Generic Functions

We get an error:

```
error[E0369]: binary operation `>` cannot be applied to type `&T`
--> src/main.rs:4:17
4 |         if item > largest {
   |                ^ ----- &T
   |                |
   |                &T
help: consider restricting type parameter `T`
1 | fn largest<T: std::cmp::PartialOrd>(list: &[T]) -> &T {
   |                                     ++++++
```

```

error[E0369]: binary operation `>` cannot be applied to type `&T`
  --> src/main.rs:4:17
4   |         if item > largest {
      |             ^         ----- &T
      |             |
      |             &T
help: consider restricting type parameter `T`
1   | fn largest<T: std::cmp::PartialOrd>(list: &[T]) -> &T {
      |             ++++++

```

- We cannot compare two `&T` to each other
- We've stated that `T` can be *any* type, regardless of if `T` is a type that cannot actually be compared
- Let's just follow the compiler's advice for now!

# Generic Functions

Once we make that change, this works!

```
fn largest<T: std::cmp::PartialOrd>(list: &[T]) -> &T {  
    let mut largest = &list[0];  
  
    for item in list {  
        if item > largest {  
            largest = item;  
        }  
    }  
  
    largest  
}
```

```
The largest number is 100  
The largest char is y
```

# Sneak Peek: Traits

```
use std::cmp::PartialOrd;

fn largest<T: PartialOrd>(list: &[T]) -> &T {
    let mut largest = &list[0];

    for item in list {
        if item > largest {
            largest = item;
        }
    }

    largest
}
```

- We'll talk about type restrictions with *traits* next week!
- For now, all you need to know is that we need the `PartialOrd` trait to enable

# Generic Structs

We can define structs to contain a generic type using the `<T>` syntax as well!

```
struct Point<T> {  
    x: T,  
    y: T,  
}  
  
fn main() {  
    let integer = Point { x: 5, y: 10 };  
    let float = Point { x: 1.0, y: 4.0 };  
}
```

# Generic Structs

Observe that this declaration defines both the `x` field and the `y` field to be of the same type.

```
struct Point<T> {  
    x: T,  
    y: T,  
}  
  
fn main() {  
    let wont_work = Point { x: 5, y: 4.0 };  
}
```





# Generic Structs

If we try to compile this, we get an error

```
error[E0308]: mismatched types
  --> src/main.rs:7:38
7  |         let wont_work = Point { x: 5, y: 4.0 };
   |                                     ^^^ expected integer,
   |                                     found floating-point number
```

# Generic Structs

If we want a struct that allows different generic fields to have different types, we need to define another generic type.

```
struct Point<T, U> {  
    x: T,  
    y: U,  
}  
  
fn main() {  
    let both_integer = Point { x: 5, y: 10 };  
    let both_float = Point { x: 1.0, y: 4.0 };  
    let integer_and_float = Point { x: 5, y: 4.0 };  
}
```

- Note that they can still be the same!

# Generic Enums

Recall the `Option<T>` type:

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

- This is a generic enum over `T` !

# Generic Enums

Enums can be generic over multiple types, just like structs.

```
enum Result<T, E> {  
    Ok(T),  
    Err(E),  
}
```

- This enum is generic over `T` and `E`, with each stored in a variant
- `Result<T, E>` is a very common type in the standard library that we will talk about next week!

# Generic Methods

Methods on structs can also be generic.

```
struct Point<T> {
    x: T,
    y: T,
}

impl<T> Point<T> {
    fn x(&self) -> &T {
        &self.x
    }
}

fn main() {
    let p = Point { x: 5, y: 10 };
    println!("p.x = {}", p.x());
}
```

# Generic Methods

```
impl<T> Point<T> {  
    fn x(&self) -> &T {  
        &self.x  
    }  
}
```

- Observe that we have to declare `T` after the `impl` as well as after `Point`
- This is to specify that we're implementing methods on the *type* `Point<T>`
- This is different from implementing methods on the *type* `Point<f32>`

# Generic `impl`

We could have made an implementation specific to `Point<f32>`:

```
impl Point<f32> {  
    fn distance_from_origin(&self) -> f32 {  
        (self.x.powi(2) + self.y.powi(2)).sqrt()  
    }  
}
```

- This code means that `Point<f32>` will have an additional `distance_from_origin` method on top of the methods defined for `Point<T>`

# Generic `impl`

Going back to the `Point<T, U>` example:

```
struct Point<T, U> {  
    x: T,  
    y: U,  
}
```

We could implement methods for when `x` is `i32`, but with no restrictions on `y`.

```
impl<U> Point<i32, U> {  
    fn get_sum_x(&self, other: Point<i32, U>) -> i32 {  
        self.x + other.x  
    }  
}
```



# Generic `impl`

However, this actually restricts the type of `other` to have the same generic type parameters `<i32, U>`.

```
impl<U> Point<i32, U> {
    fn get_sum_x(&self, other: Point<i32, U>) -> i32 {
        self.x + other.x
    }
}

fn main() {
    let p1 = Point { x: 5, y: 3.2 }; // y is f64
    let p2 = Point { x: 5, y: 4.4 }; // y is also f64
    println!("{}", p1.get_sum_x(p2))
}
```

- Note that `U` has to be the same in both `self` and `other`

# Generic `impl`

To solve this, we can make the method generic over another type:

```
impl<U> Point<i32, U> {
    fn get_sum_x<V>(&self, other: Point<i32, V>) -> i32 {
        self.x + other.x
    }
}

fn main() {
    let p1 = Point { x: 5, y: 3.2 }; // y is f64
    let p2 = Point { x: 5, y: String::new() }; // y is String
    println!("{}", p1.get_sum_x(p2))
}
```

Here's another example of a generic `impl` :

```
struct Point<X1, Y1> {
    x: X1,
    y: Y1,
}

impl<X1, Y1> Point<X1, Y1> {
    fn mixup<X2, Y2>(self, other: Point<X2, Y2>) -> Point<X1, Y2> {
        Point {
            x: self.x,
            y: other.y,
        }
    }
}

fn main() {
    let p1 = Point { x: 5, y: 10.4 };
    let p2 = Point { x: "Hello", y: 'c' };

    let p3 = p1.mixup(p2);
    println!("p3.x = {}, p3.y = {}", p3.x, p3.y);
}
```

# Performance of Generics

- The good news is that there is *zero* overhead to using generics!
  - The work is done at compile-time instead of runtime.
- Rust accomplishes this with *monomorphization*

# Monomorphization

Let's look at how this works by using the standard library's generic `Option<T>`:

```
let integer = Some(5);  
let float = Some(5.0);
```

- The compiler will identify which types `T` can take on by find all instances of `Option<T>`, in this case `i32` and `f64`
- It creates monomorphized versions of `Option` specific to those types

# Monomorphization

The compiler will generate something similar to the following:

```
enum Option_i32 {
    Some(i32),
    None,
}

enum Option_f64 {
    Some(f64),
    None,
}

fn main() {
    let integer = Option_i32::Some(5);
    let float = Option_f64::Some(5.0);
}
```

## Recap: Generics

- Generics allow us to reduce code duplication
- Monomorphization means we do not incur any runtime cost!

# Homework 4

- You'll be implementing two collection data structures:
  - `MultiSet`
    - A collection that stores unordered values and tracks multiplicity
  - `MultiMap`
    - A collection that maps keys to any number of values
- Make sure you are familiar with the API for `HashMap` and `Entry` !



## Next Lecture: Errors and Traits

Thanks for coming!

*Slides created by:*

Connor Tsui, Benjamin Owad, David Rudo,  
Jessica Ruan, Fiona Fisher, Terrance Chen

