INTRO TO RUST LANG

MODULES AND TESTING

Benjamin Owad, David Rudo, and Connor Tsui

Today: Modules and Testing

- Packages and Crates
- Modules
  - The `use` keyword
  - Module Paths and File System
- Unit Testing
- Integration Testing

Large Programs

As your programs get larger, the organization of the code becomes increasingly important.

It is generally good practice to:

- Split code into multiple folders and files
- Group related functionality
- Separate code with distinct features
- Encapsulate implementation details
- *Modularize* your program

Module System

Rust implements a number of organizational features, collectively referred to as the *module system*.

- **Packages**: A Cargo feature that lets you build, test, and share crates
- **Crates**: A tree of modules that produces a library or executable
- **Modules**: Lets you control the organization, scope, and privacy of paths
- **Paths**: A way of naming an item, such as a struct, function, or module

Packages and Crates

Crate

A *crate* is the smallest amount of code that the Rust compiler considers at a time.

- The equivalent in C/C++ is a *compilation unit*
- Running `rustc` on a single file also builds a crate
- Crates contain modules
  - Modules can be defined in other files
  - Paths allow modules to refer to other modules

Crate

There are two types of crates: binary crates and library crates.

- A binary crate can be compiled to an executable
  - Contains a `main` function
  - Examples include command-line utilities or servers
- A library crate has no `main` function, and does not compile to an executable
  - Defines functionality intended to be shared with multiple projects
- Each crate also has a file referred to as the *crate root*
  - *The Rust compiler looks at this file first, and it is also the root module of the crate (more on modules later!)*

Package

A package is a bundle of one or more crates.

- A package is defined by a `Cargo.toml` file at the root of your directory
  - `Cargo.toml` describes how to build all of the crates
- A package can contain any number of binary crates, but at most one library crate

`cargo new`

Let's walk through what happens when we create a package with `cargo new`.

```
$ cargo new my-project
     Created binary (application) `my-project` package

$ ls my-project
Cargo.toml
src

$ ls my-project/src
main.rs
```

- Creates a new package called `my-project`
- Creates a `src/main.rs` file that prints `"Hello, world!"`
- Creates a `Cargo.toml` in the root directory

`Cargo.toml`

Let's take a look inside the `Cargo.toml`.

```toml
[package]
name = "my-project"
version = "0.1.0"
edition = "2021"

[dependencies]
```

- File written in `toml`, a file format for configuration files
- Notice how there is no explicit mention of `src/main.rs`
- Cargo follows the convention that a `src/main.rs` file is the crate root of a *binary* crate
- Similarly, a `src/lib.rs` file is the crate root of a *library* crate

Example: `cargo`

Cargo is actually a Rust package that ships with installations of Rust!

- Contains the binary crate that compiles to the executable `cargo`
- Contains a library crate that the `cargo` binary depends on

Aside: Package vs Project vs Program

- "Package" is the only term of these three with a formal definition in Rust
- "Project" is a very overloaded term
  - Meaningful in the context of an *IDE*
- "Program"
  - Ask the mathematicians ¯\(ツ)/¯

Modules

Modules

Modules let us organize code within a crate for readability and easy reuse.

- Modules are collections of *items*
  - Items are functions, structs, traits, etc.
- Allows us to control the privacy of items
- Mitigates namespace collisions
- Here is a cheat sheet from the Rust Book!

Root Module

The root module is in our `main.rs` (for a binary crate) or `lib.rs` (for a library crate).

```
$ cargo new restaurant
```

src/main.rs

```rust
fn main() {
    println!("Hello, world!");
}
```

Declaring Modules

We can declare a new module with the keyword `mod`.

src/main.rs

```rust
fn main() {
    println!("Hello, World!");
}

mod kitchen {
    // `cook` is defined in the module `kitchen`
    fn cook() {
        println!("I'm cooking");
    }
}
```

Using Modules

To use items outside of a module, we must declare them as `pub`.

src/main.rs

```rust
fn main() {
    kitchen::cook();
}

mod kitchen {
    pub fn cook() { println!("I'm cooking"); }

    // Only items internal to the `kitchen` should be able to access this
    fn examine_ingredients() {}
}
```

- By default, all module items are private in Rust

Declaring Submodules

We can declare submodules inside of other modules.

src/main.rs

```rust
fn main() {
    kitchen::stove::cook();
}

mod kitchen {
    pub mod stove {
        pub fn cook() { println!("I'm cooking"); }
    }

    fn examine_ingredients() {}
}
```

- Submodules also have to be declared as `pub mod` to be accessible
- The module system is a tree, just like a file system

Modules as Files

In addition to declaring modules *within* files, creating a file named `module_name.rs` declares a corresponding module named `module_name`.

```
src
├── module_name.rs
└── main.rs
```

- Allows us to represent the module structure in the file system
- Let's try moving the `kitchen` module to its own file!

Modules as Files

src/main.rs

```rust
mod kitchen; // The compiler will look for kitchen.rs

fn main() {
    kitchen::stove::cook();
}
```

src/kitchen.rs

```rust
pub mod stove {
    pub fn cook() { println!("I'm cooking"); }
}

fn examine_ingredients() {}
```

- What about moving the `stove` submodule to its own file?

Submodules as Files

We can move the `stove` submodule into a file `src/kitchen/stove.rs` to indicate that `stove` is a submodule of `kitchen`.

src/kitchen.rs

```rust
pub mod stove; // note this still has to be `pub`

fn examine_ingredients() {}
```

src/kitchen/stove.rs

```rust
pub fn cook() {
    println!("I'm cooking");
}
```

- `main.rs` is unchanged (omitted for slide real estate)

Alternate Submodule File Naming

We could also replace `src/kitchen.rs` with `src/kitchen/mod.rs`.

src/kitchen/mod.rs

```rust
pub mod stove;

fn examine_ingredients() {}
```

src/kitchen/stove.rs

```rust
pub fn cook() {
    println!("I'm cooking");
}
```

- The only difference is in which file the `kitchen` module is defined

Alternate Submodule File Naming

In terms of Rust's module system, these two file trees are (essentially) identical.

```
src
├── kitchen
│   └── stove.rs
├── kitchen.rs
└── main.rs
```

```
src
├── kitchen
│   ├── mod.rs
│   └── stove.rs
└── main.rs
```

- This is a stylistic choice that each instructor has a very strong opinion on
  - Ask at your own peril...
- Consistency with surrounding codebase is ***always*** most important

The Module Tree, Visualized

Even with our file system changes, the module tree stays the same!

```
crate restaurant
├── mod kitchen: pub(crate)
│   ├── fn examine_ingredients: pub(self)
│   └── mod stove: pub
│       └── fn cook: pub
└── fn main: pub(crate)
```

- We can customize our file structure without changing any behavior

Module Paths

To use any item in a module, we need to know its *path*, just like a filesystem.

There are two types of paths:

- An *absolute path* is the full path starting from the crate root
- A *relative path* starts from the current module and use `self`, `super`, or an identifier in the current module
- Components of paths are separated by double colons (`::`)

Paths for Referring to Modules

You may have noticed a path from the previous sequence:

```
kitchen::stove::cook();
```

This is saying:

- In the module `kitchen`
  - In the submodule `stove`
    - Call the function `cook`
- This is a path relative to the crate root

Using Paths

src/main.rs

```rust
mod kitchen;

fn main() {
    kitchen::stove::cook();
}
```

- Not too hard to write...

Using Verbose Paths

What if we had a deeper module tree?

src/main.rs

```
fn main() {
    kitchen::stove::stovetop::burner::gas::gasknob::pot::cook();
    kitchen::stove::stovetop::burner::gas::gasknob::pot::cook();
    kitchen::stove::stovetop::burner::gas::gasknob::pot::cook();
}
```

- A lot more verbose...
  - Especially if we need to write this multiple times

The `use` Keyword

We can bring paths into scope with the `use` keyword.

src/main.rs

```rust
mod kitchen;

use kitchen::stove::stovetop::burner::gas::gasknob::pot;

fn main() {
    pot::cook();
    pot::cook();
    pot::cook();
}
```

- It is idiomatic to `use` up to the *parent* of a function, rather than the function item itself

More `use` Syntax

We can also import items from the Rust standard library (`std`).

```
use std::collections::HashMap;
use std::io::Bytes;
use std::io::Write;
```

- `HashMap` and `Bytes` are structs, and `Write` is a trait
- It is idiomatic to import structs, enums, traits, etc. directly
  - No real reason behind this besides convention

More `use` Syntax

We can combine those 2 `std::io` imports into one statement:

```
use std::collections::HashMap;
use std::io::{Bytes, Write};
```

- You could also write `use std::io::*` to bring in everything from the `std::io` module (including `Bytes` and `Write`)
  - Called the "glob operator"
  - Generally not recommended (increases compilation cost)

Aside: Binary and Library Crate Paths

In the past examples, we were using a binary crate (`src/main.rs`). All the same principles apply to using a library crate.

However, if you use *both* a binary *and* a library crate, things are slightly different.

```
src
├── kitchen
│   ├── mod.rs
│   └── stove.rs
├── lib.rs <- What happens when we add this?
└── main.rs
```

Aside: Binary and Library Crate Paths

Typically when you have both a binary and library crate in the same package, you want to call functions defined in `lib.rs` from `main.rs`.

```
src
├── kitchen
│   ├── mod.rs
│   └── stove.rs
├── lib.rs
└── main.rs (wants to call functions from lib.rs)
```

- If you have both a `main.rs` file and a `lib.rs` file, *both* are crate roots
- So how can we get items from a separate module tree?

Accessing Library from Binary

Let's try to refactor our previous example:

src/lib.rs

```rust
pub mod kitchen; // Now marked `pub`!
```

src/main.rs

```rust
fn main() {
    ???::kitchen::stove::cook();
}
```

- All files in `src/kitchen` remain unchanged
- What do we put in `???`?

Accessing Library from Binary

We treat our library crate as an *external* crate, with the same name as our package.

src/main.rs

```
fn main() {
    restaurant::kitchen::stove::cook();
}
```

- Similar to how you would treat `std` as an external crate
- We'll talk about external crates more next week!

The `super` Keyword

We can also construct relative paths that begin in the parent module with `super`.

```
crate restaurant
├── mod kitchen: pub(crate)
│   ├── fn examine_ingredients: pub(self)
│   └── mod stove: pub
│       └── fn cook: pub
└── fn main: pub(crate)
```

src/kitchen/stove.rs

```rust
pub fn cook() {
    super::examine_ingredients(); // Make sure you do this before cooking!
    println!("I'm cooking");
}
```

Privacy

```
mod kitchen: pub(crate)
├── fn examine_ingredients: pub(self)
└── mod stove: pub
    └── fn cook: pub
```

src/kitchen/stove.rs

```rust
pub fn cook() {
    super::examine_ingredients(); // Make sure you do this before cooking!
    println!("I'm cooking");
}
```

- `examine_ingredients` does not need to be public in this case
- `stove` can access anything in its parent module `kitchen`
- Privacy only applies to parent modules and above

Privacy of Types

We can also use `pub` to designate structs and enums as public.

```
pub struct Breakfast {
    pub toast: String,
    seasonal_fruit: String,
}

pub enum Appetizer {
    Soup,
    Salad,
}
```

- We can mark specific fields of structs public, allowing direct access
- If an enum is public, so are its variants!

Recap: Modules

- You can split a package into crates, and crates into modules
- You can refer to items defined in other modules with paths
- All module components are private by default, unless you mark them as `pub`

Testing

Testing

> Program testing can be a very effective way to show the presence of bugs, but it is hopelessly inadequate for showing their absence.

- Edsger W. Dijkstra, *The Humble Programmer*

Testing

Correctness of a program is complex and not easy to prove.

- Rust's type system helps with this, but it certainly cannot catch everything
- Rust includes a testing framework for this reason!

What is a Test?

Generally we want to perform at least 3 actions when running a test:

1. Set up needed data or state
2. Run the evaluated code
3. Determine if the results are as expected

Writing Tests

In Rust, a test is a function annotated with the `#[test]` attribute.

src/lib.rs

```rust
#[cfg(test)]
mod tests {
    #[test]
    fn it_works() {
        let result = 2 + 2;
        assert_eq!(result, 4);
    }
}
```

- After running `cargo new adder --lib`, this code will be in `src/lib.rs`

Writing Tests

Let's break this down.

```rust
#[test]
fn it_works() {
    let result = 2 + 2;
    assert_eq!(result, 4);
}
```

- The `#[test]` attribute indicates that this is a test function
- We set up the value `result` by adding `2 + 2`
- We use the `assert_eq!` macro to assert that `result` is correct
- We don't need to return anything, since not panicking *is* the test!

Running Tests

We run tests with `cargo test`.

```
$ cargo test
   Compiling adder v0.1.0 (file:///projects/adder)
    Finished test [unoptimized + debuginfo] target(s) in 0.57s
     Running unittests src/lib.rs (target/debug/deps/adder-92948b65e88960b4)

running 1 test
test tests::it_works ... ok

test result: ok. 1 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s

   Doc-tests adder

running 0 tests

test result: ok. 0 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s
```

Running Tests

Let's break down the output of `cargo test`.

```
running 1 test
test tests::it_works ... ok

test result: ok. 1 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s
```

- We see `test result: ok`, meaning we have passed all the tests
- In this case, only 1 test has run, and it has passed
- The `0 measured` statistic is for benchmark tests, which are currently only available in "nightly" versions of Rust

Documentation Tests

You may have seen something similar to this in your homework:

```
   Doc-tests adder

running 0 tests

test result: ok. 0 passed; 0 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s
```

- All of the code examples in documentation comments are treated as tests!
- This is useful for keeping your docs and code in sync

`#[cfg(test)]`

You may have also noticed this `#[cfg(test)]` attribute in your homework:

```
#[cfg(test)]
mod tests {
    // <-- snip -->
}
```

- This tells the compiler that this entire module should *only* be used for testing
- Effectively removes this module from the source code when compiling with `cargo build`

Writing Better Tests

Let's try and be more creative with our tests.

```rust
#[cfg(test)]
mod tests {
    #[test]
    fn exploration() {
        assert_eq!(2 + 2, 4);
    }

    #[test]
    fn another() {
        panic!("Make this test fail");
    }
}
```

Failing Tests

Let's see what we get:

```
$ cargo test
   Compiling adder v0.1.0 (file:///projects/adder)
    Finished test [unoptimized + debuginfo] target(s) in 0.72s
     Running unittests src/lib.rs (target/debug/deps/adder-92948b65e88960b4)

running 2 tests
test tests::another ... FAILED
test tests::exploration ... ok

failures:

---- tests::another stdout ----
thread 'tests::another' panicked at 'Make this test fail', src/lib.rs:10:9
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace


failures:
    tests::another

test result: FAILED. 1 passed; 1 failed; 0 ignored; 0 measured; 0 filtered out; finished in 0.00s

error: test failed, to rerun pass `--lib`
```

Failing Tests

```
failures:

---- tests::another stdout ----
thread 'tests::another' panicked at 'Make this test fail', src/lib.rs:10:9
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace


failures:
    tests::another

test result: FAILED. 1 passed; 1 failed; <-- snip -->

error: test failed, to rerun pass `--lib`
```

- Instead of `ok`, we get that the result of `tests:another` is `FAILED`

Checking Results

We can use the `assert!` macro to ensure that something is `true`.

```
#[test]
fn larger_can_hold_smaller() {
    let larger = Rectangle {
        width: 8,
        height: 7,
    };
    let smaller = Rectangle {
        width: 5,
        height: 1,
    };

    assert!(larger.can_hold(&smaller));
}
```

Testing Equality

Rust also provides a way to check equality between two values.

```rust
#[test]
fn it_adds_two() {
    assert_eq!(4, add_two(2));
}
```

Testing Equality

If `add_two(2)` somehow evaluated to `5`, we would get this output:

```
---- tests::it_adds_two stdout ----
thread 'tests::it_adds_two' panicked at 'assertion failed: `(left == right)`
  left: `4`,
 right: `5`', src/lib.rs:11:9
note: run with `RUST_BACKTRACE=1` environment variable to display a backtrace
```

- You get a nicer error message from `assert_eq!` versus using `assert!(left == right)`

Custom Error Messages

We can also write our own custom error messages in `assert!`

```
#[test]
fn greeting_contains_name() {
    let result = greeting("Carol");
    assert!(
        result.contains("Carol"),
        "Greeting did not contain name, value was `{}`",
        result
    );
}
```

`#[should_panic]`

You may have seen something similar in your homework:

```
#[test]
#[should_panic(expected = "not less than 100")]
fn greater_than_100() {
    this_better_be_less_than_100(200);
}
```

- The `#[should_panic]` attribute says that this test expects a panic!
- Adding the `expected = "..."` means we want a specific panic message

Using `Result<T, E>` in Tests

We can also use `Result` in our tests.

```
#[test]
fn it_works() -> Result<(), String> {
    if 2 + 2 == 4 {
        Ok(())
    } else {
        Err(String::from("two plus two does not equal four"))
    }
}
```

- The test will now fail if it returns `Err`
- Allows convenient usage of `?` in tests
- Note that you can't use `#[should_panic]` on tests that return a `Result`

Controlling Test Behavior

`cargo test` compiles your code in test mode and runs the resulting test binary.

- By default, it will run all tests in parallel and prevent the output (`stdout` and `stderr`) from being displayed.
- Other testing configurations are available
- *Note that you can run* `cargo test --help`*, and* `cargo test -- --help` *for help*

Running Tests in Parallel

- Suppose each of your tests all write to some shared file on disk.
  - All tests write to a file `output.txt`
- They later assert that the file still contains that data they wrote
- You probably don't want all of them to run at the same time!

Test Threads

By default, Rust will run all of the tests in parallel, on different threads.

You can use `--test-threads` to control the number of threads running the tests.

```
$ cargo test -- --test-threads=1
```

- Generally not a good idea, since the benefits of parallelism are lost

Showing Output

If you want to prevent the capturing of output, you can use `--show-output`

```
$ cargo test -- --show-output
```

- This will print the full output of every test that is run
- With 1000 tests, this might get too verbose!
- If only we could only run a subset of the tests...

Running Tests by Name

Let's say we have 1000 tests, but only one is named `one_hundred`. We can run `cargo test one_hundred` to only run that test.

```
$ cargo test one_hundred
   Compiling adder v0.1.0 (file:///projects/adder)
    Finished test [unoptimized + debuginfo] target(s) in 0.69s
     Running unittests src/lib.rs (target/debug/deps/adder-92948b65e88960b4)

running 1 test
test tests::one_hundred ... ok

test result: ok. 1 passed; 0 failed; 0 ignored; 0 measured; 999 filtered out; finished in 0.00s
```

- Notice how there are now `999 filtered out` tests, these were the tests that didn't match the name `one_hundred`

Multiple Tests by Name

`cargo` will actually find any test that matches the name you passed in.

```
$ cargo test add
   Compiling adder v0.1.0 (file:///projects/adder)
    Finished test [unoptimized + debuginfo] target(s) in 0.61s
     Running unittests src/lib.rs (target/debug/deps/adder-92948b65e88960b4)

running 2 tests
test tests::add_three_and_two ... ok
test tests::add_two_and_two ... ok

test result: ok. 2 passed; 0 failed; 0 ignored; 0 measured; 998 filtered out; finished in 0.00s
```

- If you want an exact name, use `cargo test {name} -- --exact`

Ignoring Tests

We can ignore some tests by using the `#[ignore]` attribute.

```rust
#[test]
fn it_works() {
    assert_eq!(2 + 2, 4);
}

#[test]
#[ignore]
fn expensive_test() {
    // code that takes an hour to run
}
```

- If we only want to run ignored tests, we can run `cargo test -- --ignored`
- If we want to run all tests, we can run `cargo test -- --include-ignored`

Test Organization

The Rust community thinks about tests in terms of two main categories: unit tests and integration tests.

- Unit tests test each unit of code in isolation
- Integration tests are external to your library, testing the entire system

Unit Tests

Unit tests are almost always contained within the `src` directory.

- The convention is to create a submodule named `tests` annotated with `#[cfg(test)]` for every module you want to test
- Recall that `#[cfg(test)]` attribute on items will only compile those items when running `cargo test`, and not `cargo build`
- Prevents deploying extra code in production that is only used for testing

Testing Private Functions

Rust allows you to test private functions.

```rust
// bad style for slides
pub fn add_two(a: i32) -> i32 { internal_adder(a, 2) }
fn internal_adder(a: i32, b: i32) -> i32 { a + b }

#[cfg(test)]
mod tests {
    use super::*;

    #[test]
    fn internal() {
        assert_eq!(4, internal_adder(2, 2));
    }
}
```

Integration Tests

Integration Tests use your library in the same way any other code would.

- They can only call functions that are part of your library's public API
- Useful for testing if many parts of your library work together correctly

Integration Tests

To create integration tests, we need a `tests` directory.

```
adder
├── Cargo.lock
├── Cargo.toml
├── src
│   └── lib.rs
└── tests
    └── integration_test.rs
```

- Notice how `tests` is *outside* of `src`

Integration Tests

Since we are now external to our own library, we must import everything as if it were a 3rd-party crate.

tests/integration_test.rs

```rust
use adder;

#[test]
fn it_adds_two() {
    assert_eq!(4, adder::add_two(2));
}
```

- Note that we don't need to annotate anything with `#[cfg(tests)]`
- We can now also run test files using the *name* of the file with `cargo test --test integration_test`

Submodules in Integration Tests

As you add more integration tests, you might want to make more files in the `tests` directory to help organize them.

- You can use submodules in the `tests` directory just like in the `src` directory
- You can also use the "alternate file path" method to define non-test code

Submodules in Integration Tests

Using the alternate naming convention with `common/mod.rs` tells Rust not to treat the `common` module as an integration test file.

```
├── Cargo.lock
├── Cargo.toml
├── src
│   └── lib.rs
└── tests
    ├── common
    │   └── mod.rs
    └── integration_test.rs
```

Submodules in Integration Tests

Here is an example of using `common` in an integration test:

```
└── tests
    ├── common
    │   └── mod.rs
    └── integration_test.rs
```

```rust
use adder;

mod common;

#[test]
fn it_adds_two() {
    common::setup();
    assert_eq!(4, adder::add_two(2));
}
```

Integration Tests for Binary Crates

We cannot create integration tests for a binary crate.

- Binary crates do not expose their functions
- This is why most binary crates will be paired with a library crate, even if they don't *need* to expose any functions

Recap: Testing

- Unit tests examine parts of a library in isolation and can test private implementation details
- Integration tests check that many parts of the library work together correctly
- Even though Rust can prevent some kinds of bugs, tests are still extremely important to reduce logical bugs!

Homework 6

You'll be following the Rust Book and implementing a mini version of `grep`!

- You can do this homework in <10 minutes by copying and pasting code
- We encourage you to actually read and follow the tutorial
- You will still have to add some small extra feature once you are done!
- *Remember that if you complete 4 homeworks and show up to every lecture, you pass this course!*
- *We will only grade homework 6 at the end of the semester if your grade is not already high enough*

Next Lecture: Crates, Closures, and Iterators
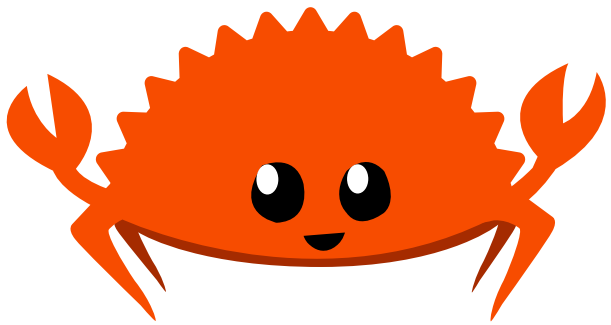
- Thanks for coming!