

A painting of a dilapidated, multi-story building with peeling paint and debris scattered around. The building has several windows and balconies. In the foreground, several large, brown crabs are scattered across the ground. The scene is set against a backdrop of a sunset or sunrise, with a large, glowing sun in the sky. The overall style is painterly and somewhat somber.

Intro to Rust Lang

Closures and

Iterators

Today: Closures and Iterators

- Closures
- Iterators
- Crate Highlights

Closures

What Is A Closure?

Closures are anonymous functions that can capture values from the scope in which they're defined.

- Known as lambdas in "lesser languages" 😊
- You can save closures in variables or pass them as arguments to other functions

Closure Syntax

```
let annotated_closure = |num: i32| -> i32 {  
    num  
};
```

- This should feel very similar to functions we've seen...
- Like normal variables, Rust can derive closure type annotations from context!

Closures Simplified

```
fn add_one_v1 (x: i32) -> i32 { x + 1 }
let add_one_v2 = |x: i32| -> i32 { x + 1 };
let add_one_v3 = |x|           { x + 1 };
let add_one_v4 = |x|           x + 1 ;

let _ = add_one_v3(3);
let _ = add_one_v4(4);
```

- `v1` is the equivalent function
- We can remove type parameters in `v3`
 - This is similar to eliding the type parameter in `let v = Vec::new()`
- For `v4`, we can remove the `{ }` since the body is only one line

How about this?

```
let example_closure = |x| x;  
  
let s = example_closure(String::from("hello"));  
let n = example_closure(5);
```



- How would we describe the type of `example_closure`?

Closure Types

```
let example_closure = |x| x;  
let s = example_closure(String::from("hello"));  
let n = example_closure(5);
```

```
error[E0308]: mismatched types
```

```
--> src/main.rs:5:29
```

```
5 |     let n = example_closure(5);  
  |                   ^- help: try using a conversion method: `.to_string()`  
  |                   |  
  |                   expected struct `String`, found integer  
  |                   arguments to this function are incorrect
```

```
note: closure parameter defined here
```

```
--> src/main.rs:2:28
```

```
2 |     let example_closure = |x| x;  
  |                           ^
```


Closure Types

```
let example_closure = |x| x;  
  
let s = example_closure(String::from("hello"));  
let n = example_closure(5);
```

- The first time we called `example_closure` with a `String`
- Rust inferred the type of `example_closure` to be `String -> String`
- Those types are now bound to the closure
 - `example_closure(5)` will not type check

Capturing References

Closures can capture values from their environment in three ways:

- Borrowing immutably
- Borrowing mutably
- Taking ownership
 - *Moving* the value into the closure

Immutable Borrowing in Closures

```
let list = vec![1, 2, 3];  
println!("Before defining closure: {:?}", list);  
  
let only_borrows = || { println!("From closure: {:?}", list); };  
  
println!("Before calling closure: {:?}", list);  
only_borrows(); // Prints "From closure: [1, 2, 3]"  
println!("After calling closure: {:?}", list);
```

```
Before defining closure: [1, 2, 3]  
Before calling closure: [1, 2, 3]  
From closure: [1, 2, 3]  
After calling closure: [1, 2, 3]
```

Immutable Borrowing in Closures

```
let list = vec![1, 2, 3];
println!("Before defining closure: {:?}", list);

let only_borrows = || { println!("From closure: {:?}", list); };

println!("Before calling closure: {:?}", list);
only_borrows(); // Prints "From closure: [1, 2, 3]"
println!("After calling closure: {:?}", list);
```

- Note how once a closure is defined, it's invoked in the same manner as a function
- Because we can have many immutable borrows, Rust allows us to print, even with the closure holding a reference

Mutable Borrowing in Closures

What happens if we mutate captured variables from inside the closure?

```
let mut list = vec![1, 2, 3];
println!("Before defining closure: {:?}", list);

let borrows_mutably = || { list.push(7); };

borrows_mutably();
println!("After calling closure: {:?}", list);
```

- This seems like it should work...



Mutable Borrowing in Closures

```
error[E0596]: cannot borrow `borrows_mutably` as mutable, as it is not declared as mutable
--> src/main.rs:7:5
5 |     let borrows_mutably = || { list.push(7); };
   |                               ---- calling `borrows_mutably` requires mutable
   |                               binding due to mutable borrow of `list`
6 |
7 |     borrows_mutably();
   |     ^^^^^^^^^^^^^^^^^ cannot borrow as mutable
```

- A closure mutating its captured state is *equivalent* to mutating *itself*
 - Calling `borrows_mutably` mutates the closure's internal state
 - We'll discuss in the next section...

Mutable Borrowing in Closures

```
help: consider changing this to be mutable
```

```
5 |     let mut borrows_mutably = || { list.push(7); };  
  |     +++
```

- As always, the compiler tells us how to fix our mistake!

Mutable Borrowing in Closures

```
let mut list = vec![1, 2, 3];
println!("Before defining closure: {:?}", list);

let mut borrows_mutably = || { list.push(7); };

borrows_mutably();
println!("After calling closure: {:?}", list);
```

```
Before defining closure: [1, 2, 3]
After calling closure: [1, 2, 3, 7]
```


Closure Borrowing Rules

```
let mut list = vec![1, 2, 3];  
  
let mut borrows_mutably = || { list.push(7); };  
  
// println!("Before calling closure: {:?}" , list); <-- Compiler error!  
borrows_mutably();  
println!("After calling closure: {:?}" , list);
```

- Note how we can't have a `println!` before invoking `borrows_mutably`
- Rust only considers the **invocation** a borrow, not the definition
 - Closures are lazy in this sense

Giving Closures Ownership

```
let mystery = {  
  let x = rand::random::<u32>();  
  |y: u32| -> u32 { x + y }  
};  
  
println!("Mystery value is {}", mystery(5));
```



Giving Closures Ownership

```
error[E0373]: closure may outlive the current block, but it borrows `x`,  
which is owned by the current block
```

```
--> src/main.rs:6:9
```

```
6 |         |y: u32| -> u32 { x + y }  
  |         ^^^^^^^^^^^^^^^^^^^^^ - `x` is borrowed here  
  |         |  
  |         may outlive borrowed value `x`
```

```
4 |     let mystery = {  
  |         ^^^^^^^
```

```
help: to force the closure to take ownership of `x`, use the `move` keyword
```

```
6 |         move |y: u32| -> u32 { x + y }
```

Giving Closures Ownership

```
let mystery = {  
    let x = rand::random::<u32>();  
    move |y: u32| -> u32 { x + y }  
// ^^^^ Add the `move` keyword!  
};  
  
println!("Mystery value is {}", mystery(5));
```

- We can `move` values into closures instead of capturing references (borrowing)
 - Note that you can't selectively `move` parameters, it's all or nothing
- `move` semantics with closures are important for thread safety!

Threads Sneak Peek

Let's briefly explore spawning a new thread with a closure.

```
fn main() {  
    let list = vec![1, 2, 3];  
    println!("Before defining closure: {:?}", list);  
  
    std::thread::spawn(move || println!("From thread: {:?}", list))  
        .join()  
        .unwrap();  
}
```

Case for `move`: Thread Safety

```
fn main() {  
    let list = vec![1, 2, 3];  
    println!("Before defining closure: {:?}", list);  
  
    std::thread::spawn(move || println!("From thread: {:?}", list))  
        .join()  
        .unwrap();  
}
```

- Why do we `move` instead of borrow?
 - Child thread's `println!` only needs a reference to `list` ...
- Parent might drop `list` before the child thread runs
 - Use after free in child thread! 💀

Handling Captured Values

- A closure body can do any of the following to a value:
 - Move a captured value out of the closure
 - Mutate a captured value
 - Neither of the above
- It could also have captured nothing to begin with!
- The properties a closure has determines the function *trait* it implements

The `Fn` traits


The Fn traits

What do you mean, function *trait*???

- Rust has 3 special traits that define the *kind* of closure we want to use
- The 3 traits are:
 - FnOnce
 - FnMut
 - Fn
- Rust auto-implements these for closures and functions

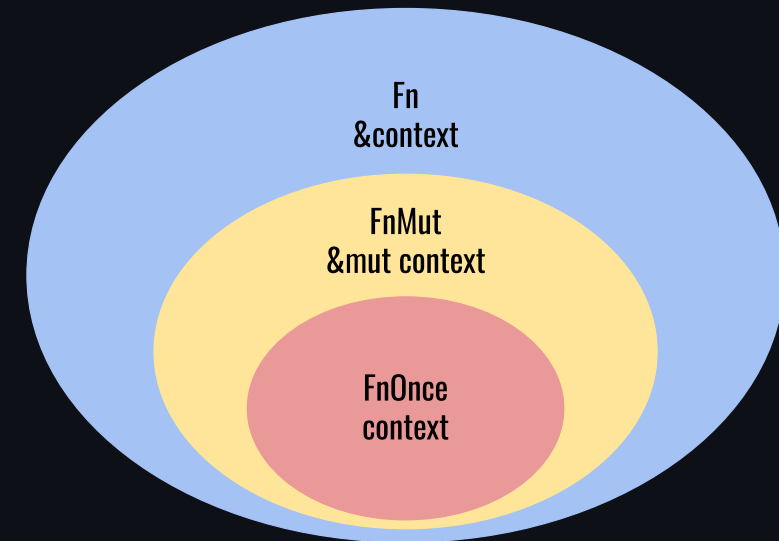
The Fn Traits: Visualized

- FnMut is also FnOnce
- Fn is also FnMut *and* FnOnce



The `Fn` traits

- `FnOnce` : Closures that can be called once
- `FnMut` : Closures that can mutate the captured values
- `Fn` : Everything else!



The `Fn` traits

- `FnOnce` : Closures that can be called once
 - All closures and functions implement this, since all closures can be called *at least once*
 - However, closures that are exclusively `FnOnce` can *only* be called once
 - e.g. A closure that moves captured values out of its body
- `FnMut` : Closures that can mutate the captured values
 - Can be called more than once
- `Fn` : Everything else!
 - Don't move values out, don't mutate, don't capture anything

FnOnce

A closure that moves captured values **out** of its body will *only* implement `FnOnce`, and not `FnMut` or `Fn`:

```
let my_str = String::from("x");  
  
// Returns `my_str`, moving it out of the closure  
let consume_and_return = move || my_str;
```

- Why can this closure only be called once?
 - It takes ownership of `my_str`, then gives ownership back to the caller
 - `my_str` is no longer accessible to our closure after it's called!
- `move` keyword specifies that the closure takes ownership when it's created, *not* when it's called

unwrap_or_else

All closures implement `FnOnce`, since all closures can be called once.

- This does not mean they can *only* be called once!
- One example is the `unwrap_or_else` method on `Result`

unwrap_or_else

`unwrap_or_else` processes the `Result<T, E>` of a function.

```
let count = |s: &str| s.len();  
  
assert_eq!(Ok(2).unwrap_or_else(count), 2);  
assert_eq!(Err("foobar").unwrap_or_else(count), 6);
```

- If successful, it unwraps the value from `Ok(T)` into `T`
- Otherwise, it takes `E` as input to the closure

unwrap_or_else

Let's look at the definition of the `unwrap_or_else` method on `Result<T>`.

```
impl<T, E> Result<T, E> {
    pub fn unwrap_or_else<F>(self, f: F) -> T
    where
        F: FnOnce(E) -> T
    {
        match self {
            Some(x) => x,
            None => f(e),
        }
    }
}
```


unwrap_or_else

First let's observe the function definition.

```
pub fn unwrap_or_else<F>(self, f: F) -> T
where
    F: FnOnce(E) -> T
// <-- snip -->
```

- This method is generic over `F`
- `F` is the type of the closure we provide when calling `unwrap_or_else`
- `F` must be able to be called once, take no arguments, and return a `T` for `Option<T>`

unwrap_or_else

Now let's observe the function body.

```
{
  match self {
    Some(x) => x,
    None => f(e),
  }
}
```

- If the `Option` is `Some`, then extract the inner value
- Otherwise, call `f` once and return the value
- Note that `f` is not *required* to only be `FnOnce` here
 - It is valid for `f` to be `FnMut` or `Fn`

FnMut

`FnMut` applies to closures that might mutate the captured values.

```
let mut x: usize = 1;
let mut add_two_to_x = || x += 2;
add_two_to_x();
```

- Use cases: Stateful operations on some shared resource
 - Imagine `x` were a score on a scoreboard
- Note that this will not compile without the second `mut`
 - `mut` signals that we are mutating our closure's environment

FnMut

Another simple example:

```
let mut base = String::from("");  
let mut build_string = |addition| base.push_str(addition);  
  
build_string("Ferris is ");  
build_string("happy!");  
  
println!("{}", base);
```

```
Ferris is happy!
```

FnMut

We can pass an `FnMut` closure as an argument to a function.

```
fn do_twice<F>(mut func: F)
where
    F: FnMut(),
{
    func();
    func();
}
```

- Would `do_twice` accept a closure that's exclusively `FnOnce`?
 - No, because we call our closure twice
- How about an `Fn` closure?
 - Yes!

Fn

The `Fn` trait is a superset of `FnOnce` and `FnMut`.

- `Fn` applies to closures that:
 - Don't move captured values out of their body
 - Don't mutate captured values
 - Don't capture anything from their environment
- Can be called more than once without mutating the environment
- Use cases:
 - Stateless operations without side effects
 - Logging, pretty printing
 - Predicates* for sorting, searching, filtering

Fn

Fn applies to closures that don't capture anything from their environment:

```
let double = |x| x * 2; // captures nothing
assert!(double(2) == 4);
```

Fn

Fn also applies to closures that don't mutate captured variables:

```
let mascot = String::from("Ferris");  
let is_mascot = |guess| guess == mascot; // `mascot` immutably borrowed  
  
assert!(is_mascot("Ferris")); // true  
assert!(!is_mascot("Ferrari")); // false
```


Fn

The `Fn` usually represents pure functions, which means it makes sense to pass them as inputs to other functions!

```
fn reduce<F, T>(reducer: F, data: &[T]) -> Option<T>
where
    F: Fn(T, T) -> T,
{
    // <-- snip -->
}
```

- We can specify the arguments and return types for `Fn`
- Functions are values!

fn vs. Fn

Rust also has function pointers, denoted `fn` (instead of `Fn`).

```
fn add_one(x: i32) -> i32 {
    x + 1
}

fn do_twice(f: fn(i32) -> i32, arg: i32) -> i32 {
    f(arg) + f(arg)
}

fn main() {
    let answer = do_twice(add_one, 5);
}
```

- `fn` is a **type*** that implements all 3 closure traits `Fn`, `FnMut`, and `FnOnce`

Recap: Closure Traits

- `Fn`, `FnMut`, `FnOnce` describe different groups of closures
 - You don't `impl` them, they apply to a closure automatically if appropriate
 - A single closure can implement one or multiple of these traits
- `FnOnce` - call at least once, environment may be consumed
- `FnMut` - call multiple times, environment may change
- `Fn` - call multiple times, environment doesn't change

Iterators

- Sorry functional haters... it's show time!

What is an Iterator?

- Iterators allow you to perform some task on a sequence of elements
- Iterators manage iterating over each item and determining termination
- Rust iterators are *lazy*
 - This means we don't pay a cost until we consume the iterator

The `Iterator` Trait

All iterators must implement the `Iterator` trait:

```
pub trait Iterator {  
    type Item;  
  
    fn next(&mut self) -> Option<Self::Item>;  
  
    // methods with default implementations elided  
}
```

- Keep generating `Some(item)`
- When the `Iterator` is finished, `None` is returned

type Item

What's going on with the `type Item` ?

```
pub trait Iterator {  
    type Item;  
  
    fn next(&mut self) -> Option<Self::Item>;  
  
    // <-- methods with default implementations elided -->  
}
```

- This is an *associated type*
- To define `Iterator` you must define the `Item` you're iterating over
- *Different from generic types!*
 - There can only be one way to iterate over something

Custom Iterator Example

Let's say we want to implement an iterator that generates the Fibonacci sequence.

```
struct Fibonacci {  
    curr: u32,  
    next: u32,  
}
```

- First need to declare the `struct` that can implement `Iterator`
- We need to store two numbers to compute the next element

Fibonacci Example

```
impl Iterator for Fibonacci {  
    type Item = u32;  
  
    fn next(&mut self) -> Option<Self::Item> {  
        let current = self.curr;  
  
        self.curr = self.next;  
        self.next = current + self.next;  
  
        // No endpoint to a Fibonacci sequence - `Some` is always returned.  
        Some(current)  
    }  
}
```

- Notice `Self::Item` is aliased to `u32`

Iterating Explicitly

```
let v1 = vec![1, 2, 3];

let mut v1_iter = v1.iter();

assert_eq!(v1_iter.next(), Some(&1));
assert_eq!(v1_iter.next(), Some(&2));
assert_eq!(v1_iter.next(), Some(&3));
assert_eq!(v1_iter.next(), None);
```

- Here we see how the `next` method is used
- Notice how `v1_iter` is mutable
 - When we call `next()` we've **consumed** that iterator element
 - The iterator's internal state has changed
 - Note that `iter()` provides immutable borrows to `v1`'s elements

Vec Iterators

```
let v1 = vec![1, 2, 3];

for val in v1.iter() {
    println!("Got: {}", val);
}

for val in &v1 {
    println!("Got: {}", val);
}
```

- These do the same thing!
 - Note that this is *not* syntax sugar
- We saw this code before in lecture 4
 - Except now we explicitly create the iterator that Rust did for us

Syntactic Sugar: `for` loops

```
let v1 = vec![1, 2, 3];

for val in v1.iter() {
    println!("Got: {}", val);
}

let mut v1_iter = v1.iter();
while let Some(val) = v1_iter.next() {
    println!("Got: {}", val);
}
```

- The `for` loop *is* syntax sugar

Iterators and Mutable Borrows

We can use the `iter_mut()` on `Vec` to get an iterator over mutable references.

```
let mut vec = vec![1, 2, 3]; // Note that we need `vec` to be mutable
let mut mutable_iter = vec.iter_mut();

while let Some(val) = mutable_iter.next() {
    *val += 1;
}

println!("{:?}", vec);
```

```
[2, 3, 4]
```

- Before we saw that `v1.iter()` gave us an iterator over `&i32`
- Now, we see that `v1.iter_mut()` gives us an iterator over `&mut i32`

More Mutable Iteration

```
let mut vec = vec![1, 2, 3];

let mut mutable_iter = vec.iter_mut();
while let Some(val) = mutable_iter.next() {
    *val += 1;
}

for val in vec.iter_mut() {
    *val += 1;
}

for val in &mut vec {
    *val += 1;
}
```

- All of these do the same thing!

Iterators and Ownership

If we want an iterator to iterate over *owned* values, we usually use `into_iter()`.

```
let mut vec = vec![1, 2, 3];

let owned_iter = vec.into_iter(); // vec is *consumed*
for val in owned_iter {
    println!("I own {}", val);
}

// owned_iter has been consumed
```

- To make an iterator that owns its values we have `into_iter()`
- `into_iter()` is the sole method on the `IntoIterator` trait
- This is what the `for` loops are actually consuming

Consuming Iterators

The standard library has many functions for using iterators.

```
let v1 = vec![1, 2, 3];  
  
let v1_iter = v1.iter();  
  
let total: i32 = v1_iter.sum(); // `.sum()` consumes `v1_iter`  
  
assert_eq!(total, 6);
```

- Most of these functions *consume* iterators


Other consuming functions

- `fn sum(self)`
- `fn max(self)`
- `fn count(self)`
- `fn map(self, f: F)`
- `fn filter(self, predicate: P)`
- `fn fold(self, init: B, f: F)`
- `fn collect(self)`
 - Coming soon...
- Many, many more!

Producing Iterators

```
let v1: Vec<i32> = vec![1, 2, 3];  
v1.iter().map(|x| x + 1);
```

- This code seems fine...



Producing Iterators

```
warning: unused `Map` that must be used  
--> src/main.rs:4:5
```

```
4 |         v1.iter().map(|x| x + 1);  
  |         ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
  |
```

```
= note: iterators are lazy and do nothing unless consumed  
= note: `[warn(unused_must_use)]` on by default
```

- Zero-cost abstractions at work
- Rust won't make us pay for our iterator until we use it
 - It will compile and warn us of unused data

Producing Iterators

```
let v2: Vec<i32> = (1..4).map(|x| x + 1).collect();  
  
println!("{:?}", v2);
```

```
[2, 3, 4]
```

- We use `collect()` to tell Rust we're done modifying our iterator and want to convert our changes to a `Vec`
- `collect` is a super common method on iterators
 - If you don't use `collect`, no computation is performed
 - `collect` is the method that executes all of the desired operations!

Filter

```
fn filter_by(list: Vec<i32>, val: i32) -> Vec<i32> {  
    list.into_iter().filter(|x| x == val).collect()  
}
```

```
--> src/main.rs:2:35  
2 |     list.into_iter().filter(|x| x == val).collect()  
  |                                     ^^ no implementation for `&i32 == i32`
```



- Some iterator functions take a reference instead of ownership
- *Note how our filter closure captures the input `val` for our filtering needs!*


Filter

```
list.into_iter().filter(|&x| x == val).collect()
```

or

```
list.into_iter().filter(|x| *x == val).collect()
```

- We either explicitly match on the reference or dereference



Chaining It Together

```
let iter = (0..100).map(|x| x*x).skip(1).filter(|y| y % 3 == 0);  
  
for x in iter.take(5) {  
    print!("{}", ", ", x);  
}
```

- You can read this as:
 - Print first 5 square numbers
 - Skipping 0
 - Only divisible by 3
- Note that `filter` doesn't need a dereference for `%`

Chaining It Together

```
let iter = (0..100).map(|x| x*x).skip(1).filter(|y| y % 3 == 0);
println!("{:?}", iter);

for x in iter.take(5) {
    print!("{}", ", ", x);
}
```

```
Filter { iter: Skip { iter: Map { iter: 0..100 }, n: 1 } }
9, 36, 81, 144, 225,
```

- Notice how `iter` is just a bunch of metadata
 - Lazy iterators on display!

Iterator Recap

- Iterators is an extremely powerful structure in Rust
- View the `std` library for more info on `Iterator` methods
- Rules regarding closures and ownership still apply
- Iterators are *lazy*
 - Remember to use `.collect()` !

Crate Highlights

rand

The standard library includes many things... but a random number generator isn't one of them*.

Here's an example of using the `rand` crate:

```
use rand::prelude::*;

let mut rng = rand::thread_rng();
let y: f64 = rng.gen(); // generates a float between 0 and 1

let mut nums: Vec<i32> = (1..100).collect();
nums.shuffle(&mut rng);
```

rand

```
use rand::prelude::*;

let mut rng = rand::thread_rng();
let y: f64 = rng.gen(); // generates a float between 0 and 1

let mut nums: Vec<i32> = (1..100).collect();
nums.shuffle(&mut rng);
```

- **rand** is the de facto crate for:
 - Generating random numbers
 - Creating probabilistic distributions
 - Providing randomness related algorithms (like vector shuffling)

clap

Often, we want our binary to take in command line arguments. A very popular argument parser used in Rust programs is `clap`.

```
use clap::Parser;

#[derive(Parser, Debug)]
#[command(version, about, long_about = None)]
struct Args {
    #[arg(short, long)]
    name: String, // Name of the person to greet

    #[arg(short, long, default_value_t = 1)]
    count: u8, // Number of times to greet
}
```

- Makes use of Rust's macro system to generate boilerplate code for us!

clap

Here's how you would use a `clap` struct called `Args` :

```
use clap::Parser;

// <-- snip -->
struct Args {
    // <-- snip -->
}

fn main() {
    let args = Args::parse(); // get-opt could never
    for _ in 0..args.count {
        println!("Hello {}!", args.name)
    }
}
```

clap

If we run the binary called `demo` :

```
$ demo --help
A simple to use, efficient, and full-featured Command Line Argument Parser

Usage: demo[EXE] [OPTIONS] --name <NAME>

Options:
  -n, --name <NAME>      Name of the person to greet
  -c, --count <COUNT>  Number of times to greet [default: 1]
  -h, --help              Print help
  -V, --version           Print version

$ demo --name Me
Hello Me!
```

- Note that `clap` is not the only 3rd-party crate option!

anyhow

Have code that can throw multiple error types that you wish was one? Use this!

```
use anyhow::Result;

fn get_cluster_info() -> Result<ClusterMap> {
    let config = std::fs::read_to_string("cluster.json"?);
    let map: ClusterMap = serde_json::from_str(&config)?;
    Ok(map)
}
```

- Both lines return different error types, but `anyhow` allows us to return both!
- Makes errors more dynamic and ergonomic

anyhow

Another example:

```
it.detach().context("Failed to detach the important thing");  
  
let content = std::fs::read(path)  
    .with_context(|| format!("Failed to read from {}", path));
```

Other `anyhow` features include:

- Downcasting to the original error types
- Attaching custom context / error messages
- More expressive custom errors

Error Handling Libraries

In addition to `anyhow`, there is also `thiserror` and `snafu`.


- `anyhow`: Use in binaries where you don't care what kind of error has occurred
- `thiserror`: Use in libraries where you *do* care what exactly happened
- `snafu`: Newer crate, combines the functionality of both!

flamegraph

Rust powered flamegraph generator with Cargo support!

With a bit of setup, you can generate this with `cargo flamegraph`.

- Can support non-Rust projects too
- Relies on `perf` or `dtrace`



Next Lecture: Ownership Revisited

Thanks for coming!

Slides created by:

Connor Tsui, Benjamin Owad, David Rudo,
Jessica Ruan, Fiona Fisher, Terrance Chen

